

Randomization of Sparse Matrix by Vector Multiplication

ABHISHEK JAIN, ISMAIL BUSTANY, and PAOLO D'ALBERTO

A sparse matrix by vector multiplication (SpMV) is simplified by the matrix non-zero elements and how we store them. There are many SpMV applications, many matrix storage formats, and thus algorithms. However, there is no optimality without considering the architecture: for example, the CPU is one among many.

By nature, randomization is resilient to counter techniques, thus suitable to avoid worst case scenarios because we tend to reduce to an average case; however, randomization does to the best case scenario the same thing it does to the worst case, it can nudge the optimal solution off. Like preconditioning, randomization is advantageous when the matrix is reused or a constant such as in the power method, Krylov's space, or convolutions for image classifications. Differently from preconditioning we do not change the values of the matrix, we randomize row and column of the matrix. We shall show that randomization is an optimization that any architecture may take advantage although in different ways. Most importantly, any developer can consider and deploy. We shall present cases where we can improve performance by 15% on AMD-based systems.

ACM Reference Format:

Abhishek Jain, Ismail Bustany, and Paolo D'Alberto. 2020. Randomization of Sparse Matrix by Vector Multiplication . 1, 1 (June 2020), 26 pages.

1 INTRODUCTION

The obvious questions are what is randomization and why would we use it? We shall provide formal definitions in the following sections, in this context, we randomly permute rows and column of a sparse matrix before a (sparse) matrix by a (dense) vector operation. We do this because randomization is the poor man's preconditioning and we do not mean it in a pejorative sense.

Preconditioning is a method to help the convergence of an iterative solution, for example a sequence of matrix by vector operations. Each iteration does a better job in searching the space and converging to a solution. In general, it means better numerical properties and well defined properties of the matrix itself. It does not mean that each iteration is faster. We want to make each iteration faster. From a mathematical and scientific point of view may seem uninteresting. From the engineering and deployment point of view is just the beginning.

There is a common thread in the scientific community how to speed it up: multi-core systems. These are composed by multi-cores processors and GPUs. The main goal is to achieve a balanced work distribution and, when applicable, minimal communication [3, 4]. When storage strategy and algorithms must be considered together then GPUs provide the work horse for the current trust and research [1]. This research is towards optimal solutions and the authors strive for a clear and complete understanding of the software-hardware relation, and usually the hardware is composed of symmetric computational units. Interestingly, the SpMV's space and time complexity, which are small, may not warrant

Authors' address: Abhishek Jain; Ismail Bustany; Paolo D'Alberto.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2020 Association for Computing Machinery.

Manuscript submitted to ACM

more performance because we end up using only one thousandth the capacity of the hardware. We may deploy efficient solutions: not necessarily faster but overall tailored for this. **FPGA or custom hardware one paragraph**

At this stage, we have too many knobs and tools to tune: algorithms, data structures, and dedicated hardware (CPU, GPUs, Custom). This is a (very) hard problem and we are not here for the solution of the inverse problem: find the best Hardware-Software solution for the one matrix by vector product. We are here to provide tools, we may say naive tools, to help understand how the structure of the matrix may affect the HW-SW solution. Randomization, or versions of it, is already used by custom hardware to re-organize the data flow to reduce communications and computation bottle necks. We come to play in this arena to show *how* to use randomization if at all.

For the readers in the field of algorithms, sparse matrix by (dense) vector is basically a sorting algorithm. Bare with us, Sorting is a method to find if an element is in a list without prior or limited knowledge of the list contents. Sorting is used to prepare the matrix and to find elements in between sparse matrices and sparse vectors. In custom architectures, sorting networks are used for routing elements of the matrix and vector to the proper functional unit. Interestingly, The best sorting algorithm is a function of the distribution of elements. If you are stuck with a sorting algorithm and the wrong distribution, randomization may change the distribution, and you do not need to talk to any HW designer

We organize our work as follows: In Section 2, we define the matrix by vector operation; in Section 3, we define what we mean for randomization. We use randomization to create a uniform distribution in Section 5 and we measure uniformity by nothing else than entropy in Section 4. We present how we drive our experiments to show the effects of randomization in Section 6. In the last sections we present a summary of the results: we present our work loads, benchmarks, in Section 7, and the complete set of measures for an AMD CPU and GPUs system in Section 8.

2 BASIC NOTATIONS

Let us start by describing the basic notations so we can clear the obvious (or not). A Sparse-matrix by vector multiplication $SpMV$ on an (semi) ring based on the operations $(+, *)$ is defined as $\mathbf{y} = \mathbb{M}\mathbf{x}$ so that $y_i = \sum_j M_{i,j} * x_j$ where $M_{i,j}=0$ are not even represented and stored. Most of the experimental results in Section 8 are based on the classic addition $(+)$ and multiplication $(*)$ in floating point precision using 64bits (i.e., double floating point precision). SpMV based on semi-ring $(min, +)$ is a short path algorithm based on an adjacent matrix of a graph, and using a Boolean algebra we can check if two nodes are connected, which is slightly simpler.

We identify a sparse matrix \mathbb{M} of size $M \times N$ as having $O(M + N)$ non-zero elements, number of non zero nnz . Thus the complexity of $\mathbb{M}\mathbf{x}$ is $O(M + N) = 2nnz$. Of course, the definition of sparsity may vary. We represent the matrix \mathbb{M} by using the Coordinate COO or and the compressed sparse row CSR ¹ format. The COO represents the non-zero of a matrix by a triplet (i, j, val) , very often there are three identical-in-size vectors for the ROW, COLUMN, and VALUE. The COO format takes $3 \times nnz$ space and two consecutive elements in the value array are not bound to be neither in the same row nor column. In fact, we know only that $VALUE[i] = M_{ROW[i], COLUMN[i]}$.

The CSR stores elements in the same row and with increasing column values consecutively. There are three arrays V , COL , and ROW . The ROW is sorted in increasing order, its size is M , and $ROW[i]$ is an index in V and COL describing where row- i starts (i.e., if row i exists). We have that $M_{i,*}$ is stored in $V[ROW[i] : ROW[i + 1]]$ and the column are at $COL[ROW[i] : ROW[i + 1]]$ and sorted increasingly. The CSR takes $2 \times nnz + M$ space and a row vector of the matrix can be found in $O(1)$.

The computation as $y_i = \sum_j M_{i,j} * x_j$ is a sequence of dot products and the CSR representation is a natural:

¹a.k.a. Compressed row storage CRS.

$$Index = ROW[i] : ROW[i + 1]$$

$$y_i = \sum_{\ell \in Index} V[\ell] * x_{COL[\ell]}$$

The matrix row is contiguous (in memory) and contiguous rows are contiguous. The access of the (dense) vector \mathbf{x} could have no pattern. The COO format could use a little preparation: For example, we can sort the array by row and add row information to achieve the same properties of CSR; however transposing a COO matrix is just a swap of the array ROW and COL. Think about matrix multiply. As today, each dot product achieves peak performance if the reads of the vector \mathbf{x} are streamlined as much as possible and so the reads of the vector V . If we have multiple cores, each could compute a sub set of the y_i and a clean data load balancing can go a long way. If we have a few functional units, we would like to have a constant stream of independent $*$ and $+$ operations but with data already in registers: that is, data pre-fetch will go a long way especially for $x_{COL[i]}$, which may have an irregular pattern.

3 RANDOMIZATION

We refer to *Randomization* as row or column permutations of the matrix \mathbb{M} (thus a permutation of \mathbf{y} and \mathbf{x}) and we choose these by a pseudo-random process. Why we want to introduce uncertainty? The sparsity of our matrix \mathbb{M} has a pattern representing the nature of the original problem; such a pattern may exploit the wrong computation for an architecture; we could break such a pattern so that the only property left is a uniform distribution (of some sort). We must avoid the worst case and we would opt for an average case instead and we could do this to a class of \mathbb{M} .

If we know the matrix \mathbb{M} and we know the architecture, preconditioning must be a better solution. Well, it is. If we run experiments long enough, we choose the best permutations for the architecture, permute \mathbb{M} , and go on testing the next. On one end, preconditioning exerts a full understanding of both the matrix (the problem) and how the final solution will be computed (architecture). This is the culminating point of knowing and we must strive to it. On the other end, the simplicity of a random permutation requires no information about the matrix, the vector, and the architecture. Such a simplicity can be exploited directly in HW. We are after an understanding when randomization is just enough: we want to let the hardware do its best with the least effort, or at least with the appearance to be effortless. Also we shall show there are different flavors of random.

Interestingly, this work stems from a sincere surprise about randomization efficacy and its application on custom SpMV. Here, we want to study this problem systematically so that to help future hardware designs. Intuitively, if we can achieve a uniform distribution of the rows of matrix \mathbb{M} we can have provable expectation of its load balancing across multiple cores. If we have a uniform distribution of accesses on \mathbf{x} we could exploit column load balancing and exploit better sorting algorithms: in practice the reading of $x_{COL[i]}$ can be reduces to a sorting and we know that different sparsity may require different algorithms. This is a lot to unpack but this translates to a better performance of the sequential algorithm without changing the algorithm or better HW utilization.

We will show that (different) randomness affects architectures and algorithms differently, making randomization a suitable optimization especially when the application and hardware are at odds, hardware is difficult to change and the matrix sparsity is simple to change. We want to show that there is a randomness hierarchy that we can distinguish as global and local; there are simple-to-find cases where the sparsity breaks randomness and the matrix has to be split into components. We want to show that this study uses common tool, open software tools and sometimes naive experiments; however, we can infer properties applicable to proprietary and custom solutions.




Fig. 1. Left: OPF 3754. Right: LP OSA 07. These are histograms where we represent normalized buckets and counts

4 ENTROPY

Patterns in sparse matrices are often visually pleasing, see Figure 1 where we present the height histogram, the width histograms and a two-dimensional histogram as heat map. We will let someone else using AI picture classification. Intuitively, we would like to express a measure of uniform distribution and here we apply the basics: *Entropy*. Given an histogram $i \in [0, M - 1]$ $h_i \in \mathbb{N}$, we define $S = \sum_{i=0}^{M-1} h_i$ and thus we have a probability distribution function $p_i = \frac{h_i}{S}$. The *information* of bin i is defined as $I(i) = -\log_2 p_i$. If we say that the stochastic variable X has PDF p_i than the entropy of X is defined as.

$$H(x) = - \sum_{i=0}^{M-1} p_i \log_2 p_i = \sum_{i=0}^{M-1} p_i I(i) = E[I_x] \quad (1)$$

The maximum entropy is when $\forall i, p_i = p = \frac{1}{M}$; that is, we are observing a uniform distributed event. There is no conceptual difference when the PDF represents a two dimensional distribution. Thus our randomization should aim at higher entropy numbers. The entropy for matrix LP OSA 07 is 8.41 and for OPF 3754 is 8.39. We use the entropy specified in the Scipy stats module. A single number is concise and satisfying. If you are pondering why they are so close contrary to their sparsity we discuss this next.

5 UNIFORM DISTRIBUTION

We know that we should **not** compare the entropy numbers of two matrices because entropy does not use any information about the order of the buckets only their probabilities. By construction, the matrices are quite different in sparsity and in shapes, however their entropy numbers are very close. Two matrices with the same number of non-zeros, spaced well enough in the proper number of bin, will have the same entropy. To appreciate their different sparsity, we should compare their entropy distributions by Jensen-Shannon measure (which is a symmetric measure, please do not use Kullback-Leibler KL divergence) [2]. Or we could use a representation of a hierarchical 2d-entropy, see Figure 2, where the entropy is split into 2x2, 4x4 and 8x8 (or fewer if the distribution is not square). We have a hierarchical entropy heat maps.

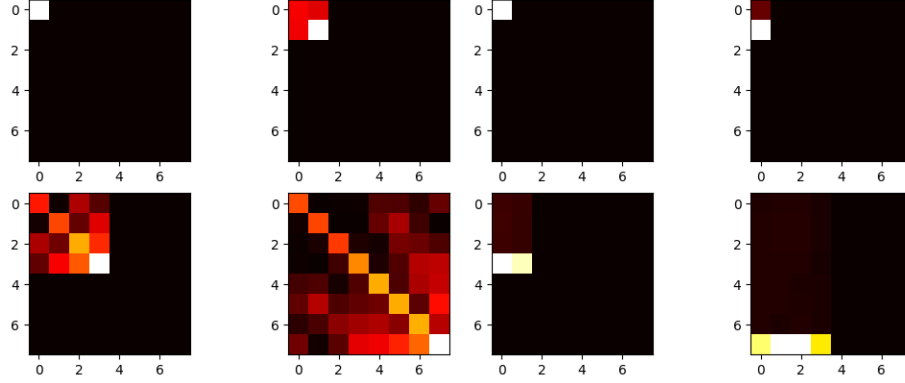


Fig. 2. Hierarchical 2D entropy for OPF 3754 (left) and LP OSA 07 (right).

We can see that a granular entropy summarizes better the nature of the matrix because it keep some spatial information. In this work, the entropy vector is used mostly for visualization purpose more than for comparison purpose. Of course, we can appreciate how the matrix LP OSA 07 has a few very heavy rows and they are clustered. This matrix will help us showing how randomization need some tips. Now we apply row and column random permutation once by row and one by column: Figure 3: OPF has now entropy 11.27 and LP 9.26. The numerical difference is significant. The good news is that for entropy, being an expectation, we can use simple techniques like bootstrap to show that the difference is significant or we have shown that Jensen-Shannon can be used and a significance level is available. What we like to see is the the hierarchical entropy heat map is becoming *more* uniform for at least one of the matrix.

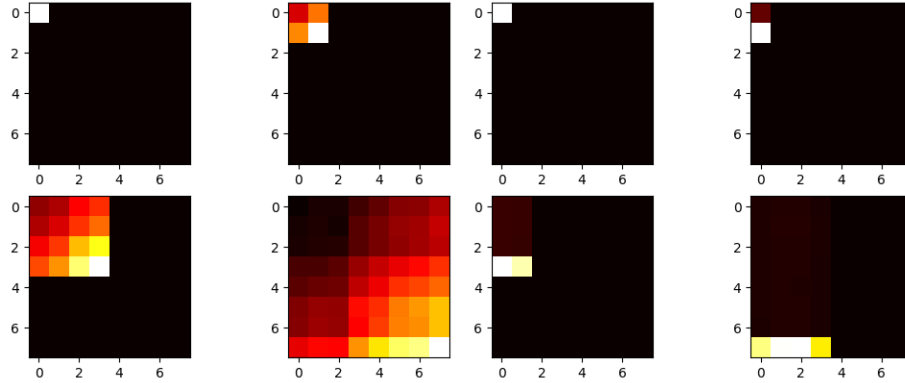


Fig. 3. Hierarchical 2D entropy after row and column random permutation for OPF 3754 (left) and LP OSA 07 (right).

In practice, permutations need some help especially for relatively large matrices. As you can see, the permutation affects locally the matrix. Of course, it depends on the implementation of the random permutation (we use numpy for this) but it is reasonable a slightly modified version of the original is still a random selection but unfortunately they seem more likely than they should. We need to compensate or help the randomization so that this current implementation does not get too lazy.

139 If we are able to identify the row and column that divide high and low density, we could use them as pivot for a
 140 shuffle like in a quick-sort algorithm. We could apply a sorting algorithm but its complexity will be the same as SpMV. We
 141 use a gradient operation to choose the element with maximum steepness, Figure 4 and 5
 142 LP achieves entropy 8.67 and 9.58 and OPF achieves 10.47 and 11.40.




Fig. 4. Hierarchical 2D entropy after height gradient based shuffle and row random permutation for OPF 3754 (left) and LP OSA 07 (right).




Fig. 5. Hierarchical 2D entropy after height and width gradient shuffle and row and column random permutation for OPF 3754 (left) and LP OSA 07 (right).

143 If the goal is to achieve a uniformly sparse matrix, it seems that we have the tools to compute and to measure such a
 144 sparsity. We admit that we do not try to find the best permutation. But our real goal is to create a work bench where
 145 randomization can be tested on different architectures and different algorithms. A randomization with a measurable
 146 uniform distribution is preferable than just random. We are interested to find out when random is enough or not enough.
 147 Also, consider that to achieve a uniform distribution, we do not need a random transformation and any permutation
 148 balancing the number of non-zero is possible, but for now not looked for.

6 MEASURING THE RANDOMIZATION EFFECTS

Whether or not this ever applied to the reader, when we have timed algorithms (i.e., measure execution time), we came to expect variation. The introduction of randomization may hide behind the ever present variance, after all these are algorithms on *small* inputs and small error can be comparable to the overall execution time. Here, we must address this concern even before describing the experiments.

First, we execute every algorithm between 1000 and 5000 times. The time of each experiment is in the seconds, providing a granularity for which we are confident the measuring time error is under control. Thus, for each experiment we provide an average execution time: we measure the time and we divide by the number of trials. Cold starts, the first iteration, are still accounted. To make the measure portable across platform we present GFLOPS, that is, Giga (10^{12}) floating operations per second: $2 * nnz$ divided by the average time in seconds.

Then we repeat the same experiment 32 times. Permutations in *numpy* Python uses a seed that is time sensitive: thus every experiment is independent from the previous. The number 32 is an old statistic trick and it is a minimum number of independent trials to approximate a normal distribution. In practice, they are not but the number is sufficient for most of the cases and it is an excellent starting point.

A short hand legend: **Reg** is the matrix without any permutation and thus is the regular; **R** stands for random Row permutation; **G-R** stands for gradient-based row shuffle and random row permutation; **G-C** stands for gradient-based column shuffle and random column permutation; **R-C** stands for random row and column permutation. This legend is used in the pictures to be concise, in the tables in the following sections, we use a verbose description. We shall clarify the gradient based approach in the experimental results section 8. Intuitively, we help the random permutation by a quick targeting of high and low volume of the histogram (and thus the matrix).

In Figure 6, We show CPU performance using COO and CSR SpMV algorithms for the matrix OPF 3754. We can see that the CSR algorithms are consistent and the Regular (i.e., the original) has always the best performance. For the COO, permutations introduce long tails, thus performance advantage.

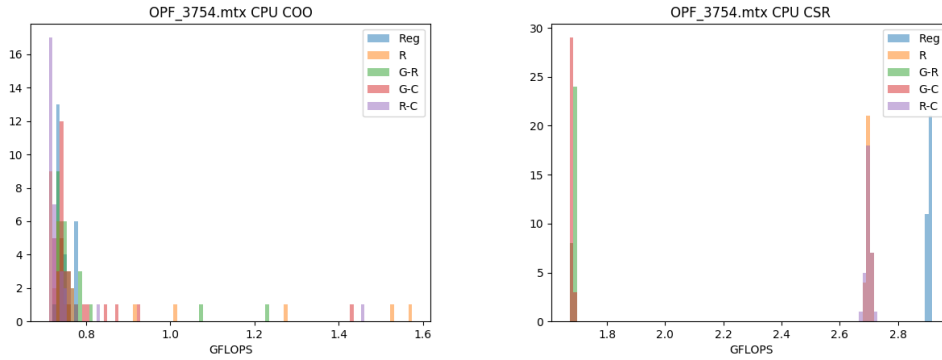


Fig. 6. CPU COO (left) and CPU CSR (left) for OPF 3754

In Figure 7, 8 and 9, randomization is harmful to the GPU implementation. The OPF 375 matrix is mostly diagonal, thus the vector \mathbf{x} is read in close quarters, randomization breaks it. If the load balance is fixed (i.e., by dividing the matrix by row and in equal row), randomization is beneficial.

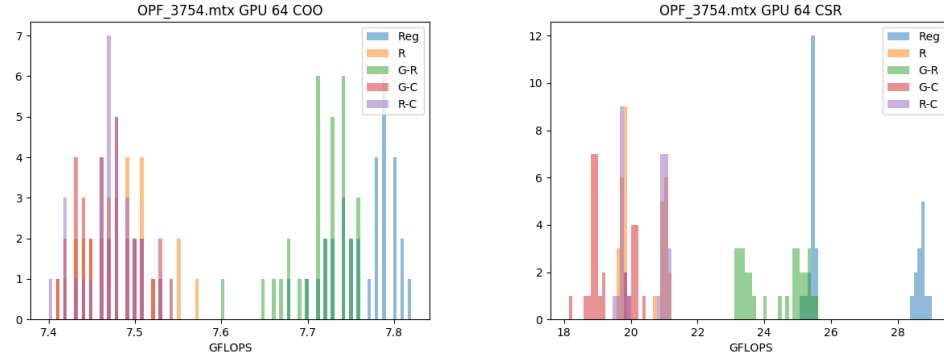


Fig. 7. Vega 20, GPU 64bits COO (left) and GPU CSR (right) for OPF 3754

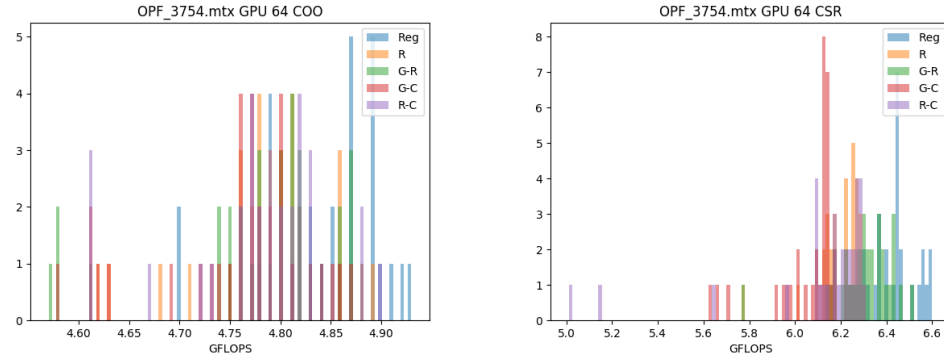


Fig. 8. Ellesmere, GPU 64bits COO (left) and GPU CSR (right) for OPF 3754

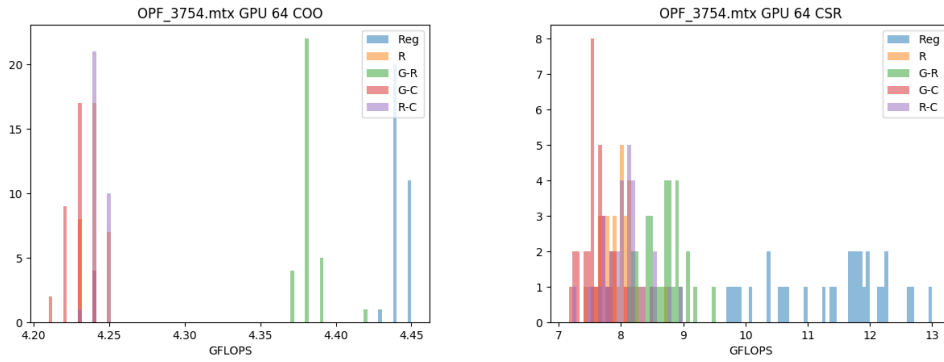


Fig. 9. Fiji, GPU 64bits COO (left) and GPU CSR (right) for OPF 3754

If we take the original matrix and split into part having the same number of rows, and execute them in parallel using different cores, we can see in Figure 10 that randomization is quite useful.




Fig. 10. Parallel CPU CSR for OPF 3754

For matrix LP OSA 07, randomization helps clearly only for CPU CSR as we show in Figure 11

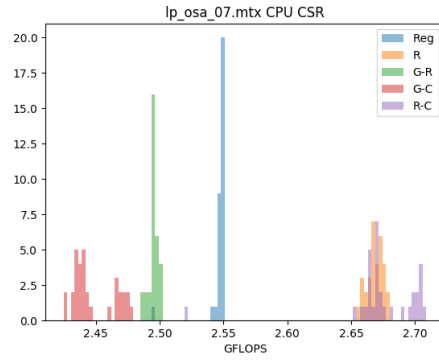


Fig. 11. CPU CSR for LP OSA 07

In Figure 12, 13, and 14, we can see that randomization is harmful but for one GPU, we can show that a single exception is possible (40% improvement).

An example, the matrix MULT DCOP 01, is where randomization is useful for the CPU, GPU, and the parallel version Figure 15, 16 - 19 and the gains can be up to 10-15%. Consider, we can achieve these improvements without any insights to the architecture, the algorithms and their relationships.

What does it mean when randomization does not work? The matrices we use in this work are not chosen randomly (pun not intended), they are the matrices that are difficult to handle in our custom SpMV engines using a combination of sorting networks and systolic arrays. If randomization does not work in our simplified work bench, will not work in our specialized architecture because the reorganization of the matrix or the input and output vector does not have the necessary parallelism, data locality, and data streaming. We need to do something else. In this case disrupting the memory pattern is not sufficient. Thus, if we cannot beat the pattern, we must exploit it, well not in this work.

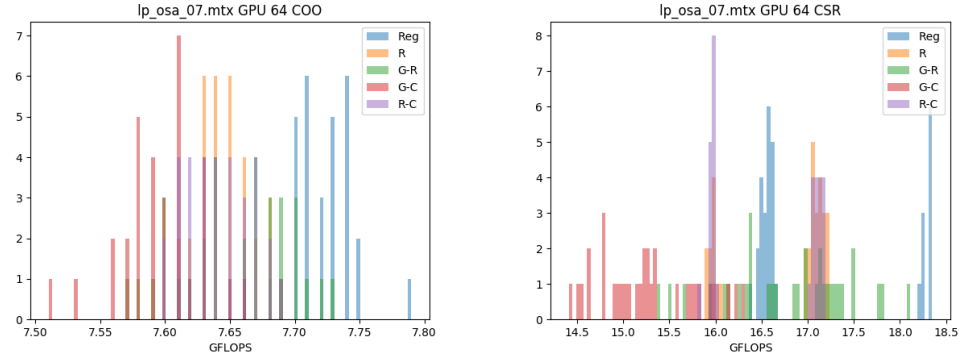


Fig. 12. Vega 20, GPU 64bits COO (left) and GPU CSR (right) for OPF 3754

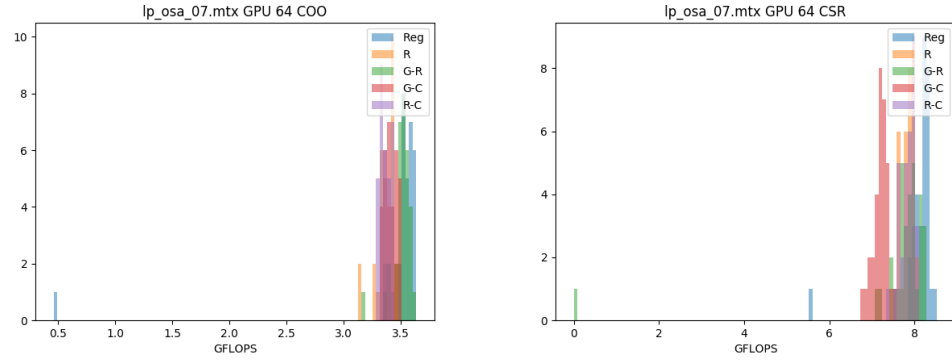


Fig. 13. Ellesmere, GPU 64bits COO (left) and GPU CSR (right) for OPF 3754

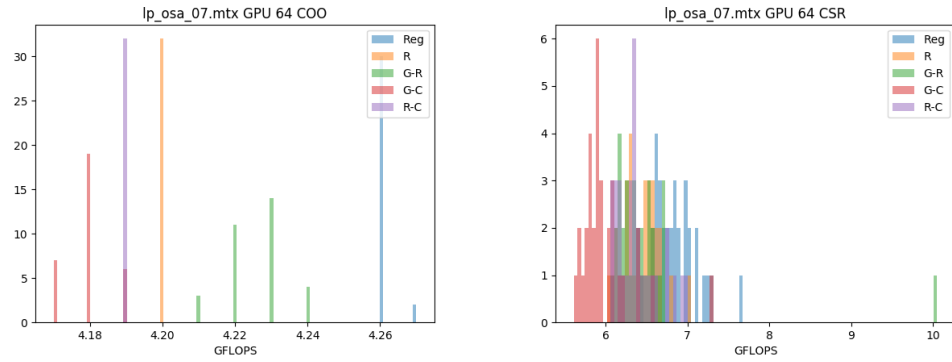


Fig. 14. Fiji, GPU 64bits COO (left) and GPU CSR (right) for OPF 3754

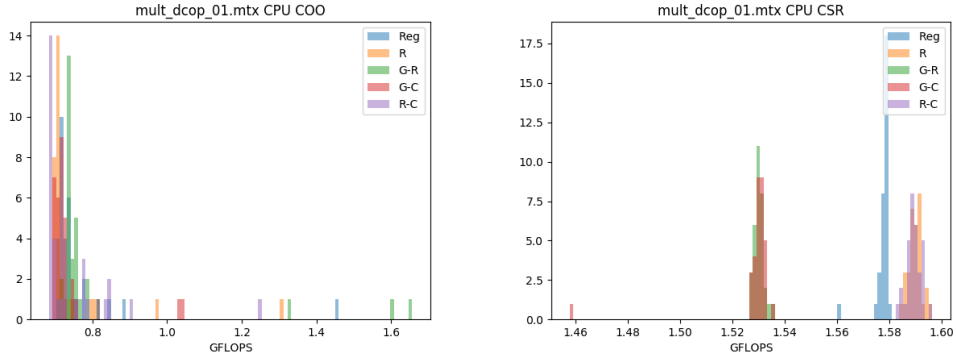


Fig. 15. CPU COO (left) and CPU CSR (right) for MULT DCOP 01

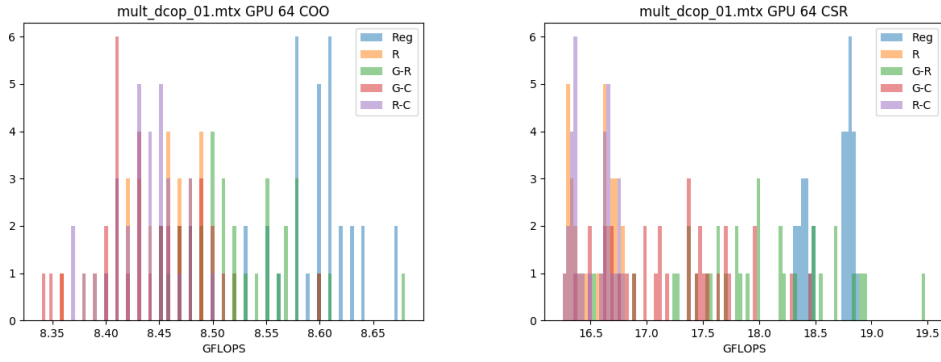


Fig. 16. Vega 20, GPU 64bits COO (left) and GPU CSR (right) for MULT DCOP 01

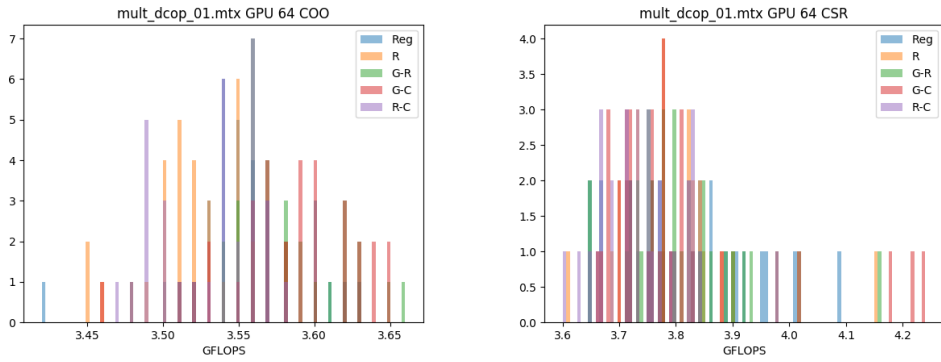


Fig. 17. Ellesmere, GPU 64bits COO (left) and GPU CSR (right) for MULT DCOP 01




Fig. 18. Fiji, GPU 64bits COO (left) and GPU CSR (right) for MULT DCOP 01




Fig. 19. Parallel CPU CSR for MULT DCOP 01

7 WORKLOADS

In the previous sections, we defined what we mean for randomization and we present our tools of tricks for the measure of the effects of randomization. Here we describe the work loads, the applications, we use to test the effects of the randomization.

7.1 Python COO and CSR algorithms

The simplicity to compute the SpMV by the code $z = A * b$ in Python is very rewarding. By change of the matrix storage format, $A = A.tocsr(); z = A * b$, we have a different algorithm. The performance exploitation is moved to the lower level. The CSR implementation is often two times faster but there are edge cases where the COO and COO with randomization can go beyond and be surprisingly better: MUL DCOP 03 is an example where COO can do well.

Intuitively, Randomization can affect the performance because the basic implementation is a sorting algorithm and it is a fixed algorithm. There are many sorting algorithms and each can be optimal for a different initial distribution. If we knew what is the sorting algorithm we could tailor the input distribution. Here we just play with it.

In Section 8, we present all the results for CPU and GPUS. Keep in mind that these problems are hard, in the sense they do not have fancy performance sheets (these architectures can achieve Tera FLOPs sustained performance for dense computations). If we go through diligently, we can see that there is a 15x performance difference between the single thread CPU and Vega 20 GPU (i.e, 3 vs 40 GFLOPS).

7.2 Parallel CSR using up to 16 cores

Python provides the concept of Pool to exploit a naive parallel computation. We notice that work given to a Pool is split accordingly to the number of elements to separate HW cores. We also noticed that the work load move from a core to another, thus not ideal. Also we notice that Pool introduce a noticeable overhead: a Pool of 1, never achieves the performance of the single thread $z = A * b$. Using Pool allows us to investigate how a naive row partitioning without counting can scale up with number of cores. We tested by splitting the rows to 1–16 cores evenly (one thread per core) and we present the performance for only the best configuration. The randomization goal is to distribute the work uniformly: a balanced work distribution avoid the unfortunate case where a single core does all the work. We are pleased by the simplicity of the benchmark and we know we can do better.

7.3 GPU COO and CSR algorithms

In this work, we use AMD GPUs and *rocSPARSE* is their current software. The software has a few glitches but overall can be used for different generation of AMD GPUs. We use the COO and CSR algorithms and we provide performance measure for double precision only. The ideas of using different GPUs: it is important to verify that the randomization can be applied independently of the HW. We are not here to compare performance across GPUs and CPUs. Often the limitation is the software, how the software can exploit the hardware or how the software will make easy to use a specific GPU. For example, the Fiji architecture is clearly superior to the Ellesmere, however the latter have better support and the system overall is more stable and user friendly.

The performance of the CSR algorithm is about two times faster than the COO. Most of the algorithms count the number of sparse elements in a row and thus they can decide the work load partition accordingly. Counting give you an edge but without changing the order of the computation there could be cases where the work load is not balanced and a little randomization could help and it does.

7.4 Randomization sometimes works

For the majority of the cases we investigated and reported in the following sections, Randomization does not work. However, there are cases where randomization does work and does work for different algorithms and architectures. If you are in the business of preconditioning, permutations are pretty cheap. If you can find a good one just consider like a preconditioning matrix, which it is.

This shows also that HW has to be more conscious, well the HW designer should, and accept that there are options at software level, at matrix level and beyond.

8 EXPERIMENTAL RESULTS

The main hardware setup is a AMD Threadripper with 16 cores. We have three Radeon GPUs: Vega 20 7nm, Pro 2xFiji, and Pro 2xEllesmere.

Vega 20 can deliver 3.5TFLOPS in double precision and it has 1TB/s HBM memory. Each Fiji provides 0.5 TFLOPS in double precision and has 512GB/s HBM, the card has two chips. The Ellesmere provides 0.3TFLOPS in double precision

238 and has 224GB/s DDR5, the card has two chips. In the performance plots presented earlier and in the following, you
 239 will notice that the performance gap between these GPUs is not so marked. We can safely state that $vega \sim 2 \times Fiji$ and
 240 $Fiji \sim 2 \times ellesmere$

241 There are 4 basic randomization formats:

- 242 • **Random Row Permutation**, we take the original matrix and permute the rows.
- 243 • **Random Row and Column Permutation**, we take the original matrix and permute the rows and the columns.
- 244 • **Gradient based row permutation**, we compute the row histogram and we compute the gradient: $h_{i+1} - h_i$.
 245 We find a single point where the gradient is maximum, this is the pivot for a shuffle like a magician would shuffle
 246 a deck of cards. Then we permute the two parts randomly.
- 247 • **Gradient based row and column permutation**, As above but also for the columns.

248 For large matrices (large number of columns and rows) a permutation tends to be a close variation of the original,
 249 still a random permutation. The gradient allows us to describe two area of the original matrix where there is a clear and
 250 de-marked density variation: for example, there are two uniform distributed sub matrices but one denser than the other.
 251 A shuffle redistribute every other sample/card to different parts and these can be permuted locally.

252 We report in the following the performance results, we introduce a * following the best performance. This is tedious
 253 to read and, we assure, to write. The code and the results are available as software repository.

9 VEGA VII AND THREADRIPPER

mult_dcop_03.mtx									
Regular									
	CPU COO	min	0.728	max	0.880	mean	0.757		
	CPU CSR	min	1.563	max	1.581	mean	1.577		
	GPU 64 COO	min	8.540	max	8.670	mean	8.619		
		CSR	min	18.320	max	18.930	mean	18.620	
	CPU PAR	min	1.170	max	1.269	mean	1.226		
	H	min	9.689	max	9.689	mean	9.689		
Row-Permute									
	CPU COO	min	0.710	max	0.845	mean	0.724		
	CPU CSR	min	1.549	max	1.597	mean	1.589		
	GPU 64 COO	min	8.360	max	8.540	mean	8.442		
		CSR	min	16.260	max	16.780	mean	16.551	
	CPU PAR	min	1.205	max	1.319	mean	1.263		
	H	min	10.737	max	10.742	mean	10.740		
Row-Gradient									
	CPU COO	min	0.706	max	1.603	mean	0.806		
	CPU CSR	min	1.493	max	1.534	mean	1.528		
	GPU 64 COO	min	8.430	max	8.610	mean	8.527		
		CSR	min	17.070	max	18.970	mean	18.115	
	CPU PAR	min	1.331	max	1.695	mean	1.513		
	H	min	10.576	max	10.585	mean	10.580		
Column-Gradient									
	CPU COO	min	0.694	max	1.632	mean	0.797		
	CPU CSR	min	1.491	max	1.534	mean	1.529		
	GPU 64 COO	min	8.350	max	8.520	mean	8.429		
		CSR	min	15.970	max	18.180	mean	17.124	
	CPU PAR	min	1.321	max	1.728	mean	1.514		
	H	min	10.826	max	10.840	mean	10.833		
Row-Column-Permute									
	CPU COO	min	0.688	max	0.757	mean	0.696		
	CPU CSR	min	1.490	max	1.595	mean	1.584		
	GPU 64 COO	min	8.380	max	8.500	mean	8.445		
		CSR	min	16.230	max	16.780	mean	16.513	
	CPU PAR	min	1.192	max	1.274	mean	1.237		
	H	min	10.737	max	10.742	mean	10.740		
mult_dcop_01.mtx									
Regular									
	CPU COO	min	0.710	max	1.453	mean	0.761		
	CPU CSR	min	1.561	max	1.581	mean	1.578		
	GPU 64 COO	min	8.520	max	8.670	mean	8.597		
		CSR	min	18.320	max	18.870	mean	18.636	
	CPU PAR	min	1.163	max	1.246	mean	1.212		
	H	min	9.689	max	9.689	mean	9.689		
Row-Permute									
	CPU COO	min	0.699	max	1.305	mean	0.745		
	CPU CSR	min	1.585	max	1.597	mean	1.590		
	GPU 64 COO	min	8.360	max	8.520	mean	8.446		
		CSR	min	16.260	max	16.780	mean	16.528	
	CPU PAR	min	1.192	max	1.298	mean	1.242		
	H	min	10.738	max	10.742	mean	10.740		
Row-Gradient									
	CPU COO	min	0.709	max	1.656	mean	0.819		
	CPU CSR	min	1.527	max	1.535	mean	1.530		
	GPU 64 COO	min	8.450	max	8.680	mean	8.527		
		CSR	min	16.520	max	19.480	mean	17.984	
	CPU PAR	min	1.280	max	1.704	mean	1.485		
	H	min	10.572	max	10.585	mean	10.581		
Column-Gradient									
	CPU COO	min	0.698	max	1.042	mean	0.737		
	CPU CSR	min	1.458	max	1.536	mean	1.528		
	GPU 64 COO	min	8.340	max	8.600	mean	8.443		
		CSR	min	16.360	max	18.450	mean	17.247	
	CPU PAR	min	1.307	max	1.712	mean	1.494		
	H	min	10.823	max	10.841	mean	10.835		

327	mult_dcop_02.mtx									
328	Regular									
329		CPU	COO	min	1.615	max*	1.677	mean	1.652	
330		CPU	CSR	min	1.539	max	1.579	mean	1.575	
331		GPU	64 COO	min	8.530	max*	8.700	mean	8.614	
332				CSR	min	18.290	max	18.890	mean	18.597
333		CPU	PAR	min	1.120	max	1.248	mean	1.211	
334		H		min	9.689	max	9.689	mean	9.689	
335	Row-Premute									
336		CPU	COO	min	0.684	max	0.780	mean	0.705	
337		CPU	CSR	min	1.558	max*	1.596	mean	1.588	
338		GPU	64 COO	min	8.360	max	8.490	mean	8.433	
339			CSR	min	16.240	max	16.750	mean	16.552	
340		CPU	PAR	min	1.182	max	1.277	mean	1.242	
341		H		min	10.737	max	10.742	mean	10.740	
342	Row-Gradient									
343		CPU	COO	min	0.704	max	1.373	mean	0.790	
344		CPU	CSR	min	1.518	max	1.535	mean	1.529	
345		GPU	64 COO	min	8.420	max	8.590	mean	8.517	
346			CSR	min	16.680	max*	19.550	mean	17.907	
347		CPU	PAR	min	1.328	max*	1.713	mean	1.484	
348		H		min	10.572	max	10.585	mean	10.581	
349	Column-Gradient									
350		CPU	COO	min	0.697	max	1.460	mean	0.742	
351		CPU	CSR	min	1.517	max	1.534	mean	1.527	
352		GPU	64 COO	min	8.330	max	8.490	mean	8.420	
353			CSR	min	16.020	max	18.390	mean	17.303	
354		CPU	PAR	min	1.321	max	1.709	mean	1.557	
355		H		min	10.823	max*	10.843	mean	10.835	
356	Row-Column-Permute									
357		CPU	COO	min	0.691	max	0.746	mean	0.698	
358		CPU	CSR	min	1.568	max	1.595	mean	1.587	
359		GPU	64 COO	min	8.350	max	8.500	mean	8.436	
360			CSR	min	16.250	max	16.780	mean	16.517	
361		CPU	PAR	min	1.187	max	1.280	mean	1.228	
362		H		min	10.739	max	10.743	mean	10.740	
363	lp_fit2d.mtx									
364	Regular									
365		CPU	COO	min	0.774	max	0.804	mean	0.793	
366		CPU	CSR	min	2.538	max	2.550	mean	2.547	
367		GPU	64 COO	min	7.060	max	7.170	mean	7.101	
368			CSR	min	15.650	max*	18.700	mean	18.031	
369		CPU	PAR	min	1.537	max	1.645	mean	1.590	
370		H		min	11.109	max	11.109	mean	11.109	
371	Row-Premute									
372		CPU	COO	min	0.740	max	0.776	mean	0.746	
373		CPU	CSR	min	3.302	max*	3.328	mean	3.317	
374		GPU	64 COO	min	7.040	max*	7.180	mean	7.098	
375			CSR	min	15.690	max	18.580	mean	16.732	
376		CPU	PAR	min	1.327	max	1.482	mean	1.422	
377		H		min	11.098	max	11.105	mean	11.101	
378	Row-Gradient									
379		CPU	COO	min	0.739	max*	2.092	mean	1.091	
380		CPU	CSR	min	2.539	max	2.546	mean	2.543	
381		GPU	64 COO	min	7.040	max	7.150	mean	7.100	
382			CSR	min	15.520	max	18.560	mean	17.547	
383		CPU	PAR	min	1.401	max	1.661	mean	1.525	
384		H		min	11.109	max				

401		CPU COO	min	0.727	max*	1.815	mean	0.892	475		GPU 64 COO	min	11.340	max*	11.860	mean	11.441
402		CPU CSR	min	2.867	max*	2.936	mean	2.917	476		CSR	min	36.010	max*	40.960	mean	38.048
403		GPU 64 COO	min	0.000	max	0.000	mean	0.000	477		CPU PAR	min	2.019	max	2.204	mean	2.130
404		CSR	min	0.000	max	0.000	mean	0.000	478		H	min	8.228	max	8.228	mean	8.228
405		CPU PAR	min	1.680	max*	1.751	mean	1.719	479	Row-Premute							
406		H	min	7.205	max	7.205	mean	7.205	480		CPU COO	min	0.718	max	0.751	mean	0.732
407	Row-Premute								481		CPU CSR	min	2.488	max	2.507	mean	2.498
408		CPU COO	min	0.678	max	1.483	mean	0.746	482		GPU 64 COO	min	10.810	max	11.090	mean	10.949
409		CPU CSR	min	2.311	max	2.326	mean	2.320	483		CSR	min	24.860	max	26.410	mean	25.527
410		GPU 64 COO	min	6.840	max*	7.270	mean	6.930	484		CPU PAR	min	1.978	max	2.290	mean	2.135
411		CSR	min	15.650	max	16.800	mean	16.233	485		H	min	11.836	max	11.840	mean	11.838
412		CPU PAR	min	1.649	max	1.730	mean	1.682	486	Row-Gradient							
413		H	min	11.026	max	11.031	mean	11.029	487		CPU COO	min	0.722	max	1.794	mean	0.769
414	Row-Gradient								488		CPU CSR	min	2.407	max	2.421	mean	2.416
415		CPU COO	min	0.708	max	1.209	mean	0.779	489		GPU 64 COO	min	11.210	max	11.480	mean	11.317
416		CPU CSR	min	1.648	max	1.735	mean	1.709	490		CSR	min	31.920	max	34.690	mean	33.246
417		GPU 64 COO	min	6.920	max	7.080	mean	7.015	491		CPU PAR	min	2.184	max*	2.302	mean	2.232
418		CSR	min	16.950	max	19.500	mean	17.794	492		H	min	10.742	max	10.757	mean	10.748
419		CPU PAR	min	1.497	max	1.743	mean	1.608	493	Column-Gradient							
420		H	min	10.298	max	10.304	mean	10.301	494		CPU COO	min	0.720	max	0.916	mean	0.742
421	Column-Gradient								495		CPU CSR	min	2.395	max	2.410	mean	2.402
422		CPU COO	min	0.709	max	1.536	mean	0.817	496		GPU 64 COO	min	10.840	max	11.070	mean	10.946
423		CPU CSR	min	1.705	max	1.753	mean	1.735	497		CSR	min	24.340	max	26.140	mean	25.393
424		GPU 64 COO	min	6.800	max	7.120	mean	6.865	498		CPU PAR	min	2.184	max	2.272	mean	2.223
425		CSR	min	15.480	max*	17.710	mean	16.470	499		H	min	11.873	max	11.882	mean	11.878
426		CPU PAR	min	1.446	max	1.718	mean	1.591	500	Row-Column-Permute							
427		H	min	10.880	max	10.886	mean	10.883	501		CPU COO	min	0.707	max	0.748	mean	0.714
428	Row-Column-Permute								502		CPU CSR	min	2.458	max	2.511	mean	2.506
429		CPU COO	min	0.670	max	1.024	mean	0.706	503		GPU 64 COO	min	10.880	max	11.070	mean	10.957
430		CPU CSR	min	2.199	max	2.340	mean	2.326	504		CSR	min	24.890	max	26.490	mean	25.642
431		GPU 64 COO	min	6.880	max	6.980	mean	6.933	505		CPU PAR	min	2.209	max	2.282	mean	2.240
432		CSR	min	15.610	max	16.900	mean	16.227	506		H	min	11.834	max*	11.840	mean	11.838
433		CPU PAR	min	1.598	max	1.668	mean	1.632	507	brainpc2.mtx							
434		H	min	11.025	max*	11.032	mean	11.029	508	Regular							
435	lp_osa_07.mtx								509		CPU COO	min	0.732	max	0.751	mean	0.744
436	Regular								510		CPU CSR	min	2.885	max*	2.916	mean	2.909
437		CPU COO	min	0.715	max	1.798	mean	0.885	511		GPU 64 COO	min	0.000	max	0.000	mean	0.000
438		CPU CSR	min	2.495	max	2.551	mean	2.547	512		CSR	min	0.000	max	0.000	mean	0.000
439		GPU 64 COO	min	7.650	max*	7.790	mean	7.718	513		CPU PAR	min	1.276	max	1.299	mean	1.286
440		CSR	min	16.390	max*	18.350	mean	17.093	514		H	min	7.478	max	7.478	mean	7.478
441		CPU PAR	min	0.963	max	1.012	mean	0.995	515	Row-Premute							
442		H	min	8.412	max	8.412	mean	8.412	516		CPU COO	min	0.727	max	0.855	mean	0.736
443	Row-Premute								517		CPU CSR	min	2.385	max	2.411	mean	2.397
444		CPU COO	min	0.720	max*	2.078	mean	1.104	518		GPU 64 COO	min	8.120	max	8.410	mean	8.206
445		CPU CSR	min	2.656	max*	2.679	mean	2.669	519		CSR	min	18.670	max	19.960	mean	19.536
446		GPU 64 COO	min	7.610	max	7.690	mean	7.647	520		CPU PAR	min	1.293	max	1.340	mean	1.314
447		CSR	min	15.910	max	17.210	mean	16.750	521		H	min	9.809	max	9.813	mean	9.811
448		CPU PAR	min	0.890	max	0.940	mean	0.918	522	Row-Gradient							
449		H	min	9.255	max	9.258	mean	9.256	523		CPU COO	min	0.696	max*	1.546	mean	0.785
450	Row-Gradient								524		CPU CSR	min	1.361	max	1.420	mean	1.411
451		CPU COO	min	0.725	max	2.078	mean	1.041	525		GPU 64 COO	min	8.190	max*	8.550	mean	8.302
452		CPU CSR	min	2.487	max	2.502	mean	2.495	526		CSR	min	18.700	max*	21.000	mean	19.890
453		GPU 64 COO	min	7.570	max	7.730	mean	7.655	527		CPU PAR	min	1.435	max	1.666	mean	1.549
454		CSR	min	15.370	max	18.100	mean	16.803	528		H	min	9.721	max	9.727	mean	9.723
455		CPU PAR	min	1.435	max	1.796	mean	1.592	529	Column-Gradient							
456		H	min	8.637	max	8.678	mean	8.672	530		CPU COO	min	0.698	max	1.467	mean	0.746
457	Column-Gradient								531		CPU CSR	min	1.377	max	1.423	mean	1.414
458		CPU COO	min	0.724	max	1.990	mean	1.000	532		GPU 64 COO	min	8.110	max	8.290	mean	8.187
459		CPU CSR	min	2.425	max	2.477	mean	2.448	533		CSR	min	18.090	max	20.190	mean	19.217
460		GPU 64 COO	min	7.510	max	7.660	mean	7.596	534		CPU PAR	min	1.345	max*	1.681	mean	1.518
461		CSR	min	14.410	max	16.290	mean	15.267	535		H	min	10.369	max*	10.372	mean	10.370
462		CPU PAR	min	1.238	max	1.774	mean	1.534	536	Row-Column-Permute							
463		H	min	9.447	max*	9.603	mean	9.576	537		CPU COO	min	0.698	max	1.390	mean	0.788
464	Row-Column-Permute								538		CPU CSR	min	2.387	max	2.410	mean	2.399
465		CPU COO	min	0.738	max	1.950	mean	1.071	539		GPU 64 COO	min	8.120	max	8.260	mean	8.191
466		CPU CSR	min	2.522	max	2.709	mean	2.675	540		CSR	min	18.530	max	19.960	mean	19.307
467		GPU 64 COO	min	7.600	max	7.690	mean	7.641	541		CPU PAR	min	1.295	max	1.347	mean	1.319
468		CSR	min	15.820	max	17.190	mean	16.572	542		H	min	9.809	max	9.813	mean	9.811
469		CPU PAR	min	0.891	max	0.944	mean	0.924	543	shermanACb.mtx							
470		H	min	9.255	max	9.258	mean	9.256	544	Regular							
471	ex19.mtx								545		CPU COO	min	0.712	max	1.201	mean	0.756
472	Regular								546		CPU CSR	min	1.558	max	1.601	mean	1.596
473		CPU COO	min	0.732	max*	1.837	mean	1.076	547		GPU 64 COO	min	7.080	max*	7.370	mean	7.184
474		CPU CSR	min	2.563	max*	2.586	mean	2.577	548		CSR	min	17.580	max*	19.480	mean	18.770

Manuscript submitted to ACM

697		CPU CSR	min	2.358	max	2.413	mean	2.392	771			CSR	min	19.960	max	21.190	mean	20.696	
698		GPU 64 COO	min	11.430	max	11.770	mean	11.549	772			CPU PAR	min	1.303	max	1.371	mean	1.345	
699		CSR	min	24.470	max	25.580	mean	24.785	773			H	min	10.059	max	10.062	mean	10.061	
700		CPU PAR	min	1.758	max	1.896	mean	1.829	774	Row-Gradient									
701		H	min	11.872	max	11.877	mean	11.875	775			CPU COO	min	0.723	max	0.984	mean	0.753	
702	Row-Gradient								776			CPU CSR	min	1.781	max	1.809	mean	1.803	
703		CPU COO	min	0.716	max	0.775	mean	0.739	777			GPU 64 COO	min	9.380	max	9.660	mean	9.464	
704		CPU CSR	min	1.651	max	1.689	mean	1.675	778			CSR	min	15.770	max	19.090	mean	18.037	
705		GPU 64 COO	min	12.100	max	12.410	mean	12.205	779			CPU PAR	min	1.775	max	1.924	mean	1.868	
706		CSR	min	31.670	max	34.910	mean	33.370	780			H	min	10.205	max	10.233	mean	10.219	
707		CPU PAR	min	2.079	max	2.286	mean	2.207	781	Column-Gradient									
708		H	min	11.111	max	11.116	mean	11.113	782			CPU COO	min	0.715	max	0.926	mean	0.757	
709	Column-Gradient								783			CPU CSR	min	1.729	max	1.802	mean	1.791	
710		CPU COO	min	0.715	max	1.021	mean	0.743	784			GPU 64 COO	min	9.080	max	9.270	mean	9.158	
711		CPU CSR	min	1.655	max	1.674	mean	1.666	785			CSR	min	13.980	max	15.780	mean	14.938	
712		GPU 64 COO	min	11.340	max	11.560	mean	11.463	786			CPU PAR	min	1.751	max	1.906	mean	1.846	
713		CSR	min	23.770	max	25.470	mean	24.489	787			H	min	11.213	max	11.232	mean	11.222	
714		CPU PAR	min	2.056	max	2.172	mean	2.118	788	Row-Column-Permute									
715		H	min	12.040	max	12.047	mean	12.043	789			CPU COO	min	0.732	max	1.598	mean	0.785	
716	Row-Column-Permute								790			CPU CSR	min	2.594	max	2.602	mean	2.599	
717		CPU COO	min	0.677	max	0.785	mean	0.687	791			GPU 64 COO	min	9.340	max	9.460	mean	9.394	
718		CPU CSR	min	2.325	max	2.434	mean	2.369	792			CSR	min	19.950	max	21.500	mean	20.544	
719		GPU 64 COO	min	11.450	max	11.650	mean	11.538	793			CPU PAR	min	1.326	max	1.374	mean	1.354	
720		CSR	min	24.330	max	25.560	mean	25.008	794			H	min	10.059	max	10.062	mean	10.061	
721		CPU PAR	min	1.631	max	1.776	mean	1.709	795	mhd4800a.mtx									
722		H	min	11.873	max	11.877	mean	11.875	796	Regular									
723	OPF_3754.mtx								797			CPU COO	min	0.759	max	0.795	mean	0.780	
724	Regular								798			CPU CSR	min	2.479	max	2.565	mean	2.557	
725		CPU COO	min	0.726	max	0.774	mean	0.747	799			GPU 64 COO	min	5.490	max	5.650	mean	5.552	
726		CPU CSR	min	2.898	max	2.919	mean	2.908	800			CSR	min	16.700	max	19.460	mean	18.004	
727		GPU 64 COO	min	7.680	max	7.820	mean	7.766	801			CPU PAR	min	1.456	max	1.523	mean	1.492	
728		CSR	min	25.070	max	29.030	mean	26.756	802			H	min	7.132	max	7.132	mean	7.132	
729		CPU PAR	min	1.437	max	1.508	mean	1.471	803	Row-Permute			CPU COO	min	0.695	max	0.943	mean	0.726
730		H	min	8.393	max	8.393	mean	8.393	804			CPU CSR	min	2.480	max	2.488	mean	2.485	
731	Row-Permute								805			GPU 64 COO	min	5.410	max	5.490	mean	5.453	
732		CPU COO	min	0.714	max	1.574	mean	0.817	806			CSR	min	15.700	max	17.520	mean	16.678	
733		CPU CSR	min	2.686	max	2.711	mean	2.699	807			CPU PAR	min	1.422	max	1.514	mean	1.474	
734		GPU 64 COO	min	7.410	max	7.570	mean	7.484	808			H	min	10.959	max	10.966	mean	10.963	
735		CSR	min	19.600	max	21.190	mean	20.307	809	Row-Gradient			CPU COO	min	0.723	max	2.029	mean	0.990
736		CPU PAR	min	1.443	max	1.505	mean	1.469	810			CPU CSR	min	2.411	max	2.427	mean	2.421	
737		H	min	11.267	max	11.272	mean	11.269	811			GPU 64 COO	min	5.490	max	5.560	mean	5.534	
738	Row-Gradient								812			CSR	min	16.350	max	19.560	mean	17.784	
739		CPU COO	min	0.723	max	1.232	mean	0.775	813			CPU PAR	min	1.441	max	1.509	mean	1.477	
740		CPU CSR	min	1.672	max	1.691	mean	1.685	814			H	min	9.512	max	9.526	mean	9.520	
741		GPU 64 COO	min	7.600	max	7.760	mean	7.716	815	Column-Gradient			CPU COO	min	0.721	max	1.802	mean	0.871
742		CSR	min	23.160	max	25.590	mean	24.304	816			CPU CSR	min	2.393	max	2.408	mean	2.404	
743		CPU PAR	min	1.675	max	1.736	mean	1.703	817			GPU 64 COO	min	5.410	max	5.480	mean	5.453	
744		H	min	10.463	max	10.472	mean	10.468	818			CSR	min	15.680	max	17.870	mean	16.540	
745	Column-Gradient								819			CPU PAR	min	1.429	max	1.488	mean	1.468	
746		CPU COO	min	0.726	max	1.431	mean	0.778	820			H	min	10.931	max	10.945	mean	10.938	
747		CPU CSR	min	1.671	max	1.685	mean	1.679	821	Row-Column-Permute			CPU COO	min	0.728	max	1.646	mean	1.037
748		GPU 64 COO	min	7.410	max	7.530	mean	7.467	822			CPU CSR	min	2.472	max	2.488	mean	2.480	
749		CSR	min	18.140	max	20.350	mean	19.315	823			GPU 64 COO	min	5.410	max	5.480	mean	5.449	
750		CPU PAR	min	1.650	max	1.736	mean	1.699	824			CSR	min	15.760	max	17.560	mean	16.654	
751		H	min	11.393	max	11.401	mean	11.397	825	Row-Column-Permute			CPU PAR	min	1.428	max	1.513	mean	1.474
752	Row-Column-Permute								826			H	min	10.959	max	10.967	mean	10.963	
753		CPU COO	min	0.711	max	1.458	mean	0.751	827			CPU COO	min	0.737	max	1.977	mean	1.431	
754		CPU CSR	min	2.678	max	2.717	mean	2.700	828			CPU CSR	min	2.674	max	2.688	mean	2.681	
755		GPU 64 COO	min	7.400	max	7.540	mean	7.471	829			GPU 64 COO	min	5.900	max	6.000	mean	5.954	
756		CSR	min	19.560	max	21.150	mean	20.453	830			CSR	min	13.650	max	15.410	mean	14.657	
757		CPU PAR	min	1.440	max	1.499	mean	1.467	831	gen4.mtx			CPU PAR	min	1.468	max	1.521	mean	1.491
758		H	min	11.266	max	11.272	mean	11.269	832	Regular			H	min	9.234	max	9.234	mean	9.234
759	c-47.mtx								833			CPU COO	min	0.740	max	2.048	mean	1.121	
760	Regular								834			CPU CSR	min	2.777	max	2.798	mean	2.790	
761		CPU COO	min	0.754	max	1.829	mean	1.204	835			GPU 64 COO	min	5.900	max	6.000	mean	5.954	
762		CPU CSR	min	2.610	max	2.624	mean	2.618	836			CSR	min	13.650	max	15.410	mean	14.657	
763		GPU 64 COO	min	9.530	max	9.870	mean	9.640	837			CPU PAR	min	1.468	max	1.521	mean	1.491	
764		CSR	min	23.990	max	25.910	mean	24.992	838			H	min	9.234	max	9.234	mean	9.234	
765		CPU PAR	min	1.311	max	1.380	mean	1.357	839	Row-Permute			CPU COO	min	0.740	max	2.048	mean	1.121
766		H	min	8.364	max	8.364	mean	8.364	840			CPU CSR	min	2.777	max	2.798	mean	2.790	
767	Row-Permute								841			GPU 64 COO	min	5.910	max	5.970	mean	5.944	
768		CPU COO	min	0.740	max	0.885	mean	0.755	842			CSR	min	13.700	max	15.370	mean	14.541	
769		CPU CSR	min	2.574	max	2.611	mean	2.597	843			CPU PAR	min	1.468	max	1.546	mean	1.502	
770		GPU 64 COO	min	9.320	max	9.510	mean	9.397	844										

845		H	min	10.250	max	10.255	mean	10.252	919		CPU COO	min	0.735	max	1.806	mean	0.878
846	Row-Gradient								920		CPU CSR	min	2.706	max	2.744	mean	2.726
847		CPU COO	min	0.740	max	1.790	mean	0.994	921		GPU 64 COO	min	6.390	max	6.500	mean	6.433
848		CPU CSR	min	2.663	max	2.682	mean	2.674	922		CSR	min	19.780	max	22.870	mean	20.936
849		GPU 64 COO	min	5.890	max*	6.160	mean	5.946	923		CPU PAR	min	1.710	max	1.865	mean	1.785
850		CSR	min	13.780	max*	17.520	mean	15.601	924		H	min	10.251	max	10.267	mean	10.257
851		CPU PAR	min	1.479	max*	1.619	mean	1.569	925	Column-Gradient							
852		H	min	9.939	max	9.955	mean	9.948	926		CPU COO	min	0.728	max	1.792	mean	0.986
853	Column-Gradient								927		CPU CSR	min	2.521	max	2.720	mean	2.703
854		CPU COO	min	0.743	max	1.991	mean	0.981	928		GPU 64 COO	min	6.280	max	6.370	mean	6.327
855		CPU CSR	min	2.620	max	2.654	mean	2.646	929		CSR	min	18.000	max	19.720	mean	19.040
856		GPU 64 COO	min	5.840	max	5.910	mean	5.885	930		CPU PAR	min	1.649	max	1.741	mean	1.702
857		CSR	min	13.130	max	17.040	mean	15.008	931		H	min	11.113	max	11.121	mean	11.117
858		CPU PAR	min	1.477	max	1.607	mean	1.559	932	Row-Column-Permute							
859		H	min	10.858	max*	10.876	mean	10.864	933		CPU COO	min	0.714	max	1.525	mean	0.957
860	Row-Column-Permute								934		CPU CSR	min	2.876	max	2.892	mean	2.884
861		CPU COO	min	0.742	max	2.010	mean	1.124	935		GPU 64 COO	min	6.280	max	6.370	mean	6.322
862		CPU CSR	min	2.789	max*	2.800	mean	2.795	936		CSR	min	17.960	max	19.670	mean	18.670
863		GPU 64 COO	min	5.900	max	5.980	mean	5.941	937		CPU PAR	min	1.667	max	1.754	mean	1.710
864		CSR	min	13.640	max	15.410	mean	14.556	938		H	min	11.162	max*	11.168	mean	11.165
865		CPU PAR	min	1.462	max	1.540	mean	1.504	939	TSOPF_RS_b39_c7.mtx							
866		H	min	10.250	max	10.253	mean	10.252	940	Regular							
867	Maragal_6.mtx								941		CPU COO	min	0.771	max	0.793	mean	0.780
868	Regular								942		CPU CSR	min	3.219	max*	3.232	mean	3.227
869		CPU COO	min	0.725	max	0.741	mean	0.729	943		GPU 64 COO	min	11.070	max*	11.200	mean	11.142
870		CPU CSR	min	2.345	max	2.409	mean	2.372	944		CSR	min	37.050	max*	42.100	mean	39.040
871		GPU 64 COO	min	18.200	max	18.770	mean	18.357	945		CPU PAR	min	1.910	max	2.027	mean	1.982
872		CSR	min	38.310	max*	40.240	mean	39.477	946		H	min	7.304	max	7.304	mean	7.304
873		CPU PAR	min	0.789	max	0.813	mean	0.797	947	Row-Premute							
874		H	min	9.930	max	9.930	mean	9.930	948		CPU COO	min	0.701	max	0.722	mean	0.707
875	Row-Premute								949		CPU CSR	min	2.931	max	2.952	mean	2.942
876		CPU COO	min	0.709	max	0.779	mean	0.715	950		GPU 64 COO	min	10.860	max	11.030	mean	10.928
877		CPU CSR	min	2.675	max	2.715	mean	2.696	951		CSR	min	28.730	max	30.880	mean	29.483
878		GPU 64 COO	min	17.810	max	18.030	mean	17.935	952		CPU PAR	min	1.760	max	1.922	mean	1.851
879		CSR	min	29.650	max	30.580	mean	30.109	953		H	min	10.537	max	10.541	mean	10.539
880		CPU PAR	min	0.857	max	0.940	mean	0.904	954	Row-Gradient							
881		H	min	10.777	max	10.779	mean	10.778	955		CPU COO	min	0.747	max	0.808	mean	0.757
882	Row-Gradient								956		CPU CSR	min	2.606	max	2.648	mean	2.624
883		CPU COO	min	0.710	max*	1.566	mean	0.755	957		GPU 64 COO	min	10.850	max	11.120	mean	10.999
884		CPU CSR	min	2.042	max	2.159	mean	2.120	958		CSR	min	33.910	max	37.600	mean	35.909
885		GPU 64 COO	min	18.460	max*	18.960	mean	18.665	959		CPU PAR	min	2.154	max*	2.245	mean	2.203
886		CSR	min	25.650	max	27.330	mean	26.549	960		H	min	9.636	max	9.646	mean	9.642
887		CPU PAR	min	2.257	max	2.612	mean	2.416	961	Column-Gradient							
888		H	min	11.251	max	11.301	mean	11.285	962		CPU COO	min	0.718	max*	1.693	mean	0.802
889	Column-Gradient								963		CPU CSR	min	2.502	max	2.585	mean	2.547
890		CPU COO	min	0.711	max	0.743	mean	0.725	964		GPU 64 COO	min	10.700	max	10.990	mean	10.804
891		CPU CSR	min	2.036	max	2.161	mean	2.110	965		CSR	min	27.230	max	29.380	mean	28.488
892		GPU 64 COO	min	17.840	max	18.860	mean	18.149	966		CPU PAR	min	2.128	max	2.227	mean	2.172
893		CSR	min	19.410	max	20.690	mean	20.066	967		H	min	11.131	max*	11.222	mean	11.208
894		CPU PAR	min	2.174	max*	2.546	mean	2.349	968	Row-Column-Permute							
895		H	min	12.011	max*	12.072	mean	12.052	969		CPU COO	min	0.709	max	0.726	mean	0.716
896	Row-Column-Permute								970		CPU CSR	min	2.917	max	2.958	mean	2.940
897		CPU COO	min	0.712	max	0.971	mean	0.737	971		GPU 64 COO	min	10.840	max	11.030	mean	10.930
898		CPU CSR	min	2.732	max*	2.751	mean	2.743	972		CSR	min	28.780	max	30.810	mean	29.578
899		GPU 64 COO	min	17.720	max	18.070	mean	17.911	973		CPU PAR	min	1.757	max	1.834	mean	1.792
900		CSR	min	29.600	max	30.500	mean	29.961	974		H	min	10.537	max	10.540	mean	10.539
901		CPU PAR	min	0.827	max	0.954	mean	0.913									
902		H	min	10.776	max	10.778	mean	10.777									
903	aft01.mtx								975	10 ELLESMERE							
904	Regular								976	aft01.mtx							
905		CPU COO	min	0.735	max*	2.079	mean	1.069	977	Regular							
906		CPU CSR	min	3.132	max*	3.154	mean	3.145	978		GPU 64 COO	min	4.080	max*	4.280	mean	4.186
907		GPU 64 COO	min	6.390	max*	6.610	mean	6.457	979		CSR	min	9.660	max*	12.660	mean	11.485
908		CSR	min	19.990	max*	23.250	mean	21.820	980		H	min	7.811	max	7.811	mean	7.811
909		CPU PAR	min	1.746	max*	1.865	mean	1.812	981	Row-Premute							
910		H	min	7.811	max	7.811	mean	7.811	982		GPU 64 COO	min	3.860	max	4.090	mean	4.001
911	Row-Premute								983		CSR	min	9.520	max	10.340	mean	9.936
912		CPU COO	min	0.714	max	1.648	mean	0.840	984		H	min	11.161	max	11.167	mean	11.165
913		CPU CSR	min	2.864	max	2.892	mean	2.883	985	Row-Gradient							
914		GPU 64 COO	min	6.280	max	6.380	mean	6.329	986		GPU 64 COO	min	4.010	max	4.240	mean	4.135
915		CSR	min	17.980	max	19.700	mean	19.105	987		CSR	min	5.890	max	11.350	mean	6.882
916		CPU PAR	min	1.729	max	1.850	mean	1.782	988		H	min	10.246	max	10.262	mean	10.256
917		H	min	11.162	max	11.168	mean	11.165	989	Column-Gradient							
918	Row-Gradient																

990		GPU 64 COO min 3.850 max 4.100 mean 4.012	1064		H min 7.380 max 7.380 mean 7.380
991		CSR min 5.460 max 8.790 mean 6.005	1065	Row-Premute	
992		H min 11.112 max 11.122 mean 11.117	1066		GPU 64 COO min 4.820 max 4.940 mean 4.859
993	Row-Column-Permute		1067		CSR min 5.080 max 6.520 mean 6.342
994		GPU 64 COO min 3.850 max 4.080 mean 3.990	1068	Row-Gradient	H min 10.042 max 10.047 mean 10.044
995		CSR min 5.420 max 6.760 mean 5.977	1069		GPU 64 COO min 4.810 max* 4.940 mean 4.876
996		H min 11.162 max*11.169 mean 11.165	1070		CSR min 6.100 max* 6.560 mean 6.307
997	blowey.mat		1071		H min 9.681 max 9.704 mean 9.694
998	Regular		1072	Column-Gradient	
999		GPU 64 COO min 0.000 max 0.000 mean 0.000	1073		GPU 64 COO min 4.810 max 4.930 mean 4.869
1000		CSR min 0.000 max 0.000 mean 0.000	1074		CSR min 4.820 max 6.460 mean 6.208
1001		H min 7.205 max 7.205 mean 7.205	1075		H min 10.554 max*10.661 mean 10.638
1002	Row-Premute		1076	Row-Column-Permute	
1003		GPU 64 COO min 3.800 max 3.940 mean 3.875	1077		GPU 64 COO min 4.810 max 4.940 mean 4.864
1004		CSR min 3.710 max 4.570 mean 4.399	1078		CSR min 5.930 max 6.520 mean 6.379
1005		H min 11.025 max 11.031 mean 11.028	1079		H min 10.041 max 10.047 mean 10.044
1006	Row-Gradient		1080		
1007		GPU 64 COO min 3.800 max* 4.120 mean 3.962	1081	cvxqp3.mat	
1008		CSR min 4.340 max* 4.670 mean 4.546	1082	Regular	
1009		H min 10.296 max 10.307 mean 10.300	1083		GPU 64 COO min 3.350 max* 3.590 mean 3.483
1010	Column-Gradient		1084		CSR min 5.430 max* 9.260 mean 8.333
1011		GPU 64 COO min 3.880 max 4.100 mean 3.978	1085		H min 8.646 max 8.646 mean 8.646
1012		CSR min 4.240 max 4.570 mean 4.412	1086	Row-Premute	
1013		H min 10.881 max 10.886 mean 10.883	1087		GPU 64 COO min 3.230 max 3.480 mean 3.371
1014	Row-Column-Permute		1088		CSR min 7.560 max 8.220 mean 7.900
1015		GPU 64 COO min 3.800 max 3.980 mean 3.885	1089	Row-Gradient	H min 11.027 max 11.033 mean 11.030
1016		CSR min 4.130 max 4.540 mean 4.399	1090		GPU 64 COO min 3.240 max 3.510 mean 3.396
1017		H min 11.025 max*11.033 mean 11.029	1091		CSR min 6.990 max 7.890 mean 7.574
1018	brainpc2.mat		1092		H min 11.060 max 11.069 mean 11.064
1019	Regular		1093	Column-Gradient	
1020		GPU 64 COO min 0.000 max 0.000 mean 0.000	1094		GPU 64 COO min 3.240 max 3.480 mean 3.374
1021		CSR min 0.000 max 0.000 mean 0.000	1095		CSR min 6.980 max 7.900 mean 7.557
1022		H min 7.478 max 7.478 mean 7.478	1096		H min 11.126 max*11.134 mean 11.130
1023	Row-Premute		1097	Row-Column-Permute	
1024		GPU 64 COO min 3.840 max* 6.750 mean 4.110	1098		GPU 64 COO min 3.110 max 3.470 mean 3.365
1025		CSR min 4.260 max* 4.580 mean 4.437	1099		CSR min 4.810 max 8.210 mean 7.742
1026		H min 9.809 max 9.813 mean 9.811	1100		H min 11.026 max 11.032 mean 11.030
1027	Row-Gradient		1101		
1028		GPU 64 COO min 0.640 max 4.030 mean 3.864	1102	ex19.mat	
1029		CSR min 4.270 max 4.470 mean 4.383	1103	Regular	
1030		H min 9.722 max 9.727 mean 9.724	1104		GPU 64 COO min 2.450 max* 2.610 mean 2.564
1031	Column-Gradient		1105		CSR min 4.490 max 4.760 mean 4.714
1032		GPU 64 COO min 0.640 max 4.070 mean 3.898	1106		H min 8.228 max 8.228 mean 8.228
1033		CSR min 4.230 max 4.580 mean 4.386	1107	Row-Premute	
1034		H min 10.368 max*10.372 mean 10.370	1108		GPU 64 COO min 2.000 max 2.040 mean 2.021
1035	Row-Column-Permute		1109		CSR min 4.640 max 4.780 mean 4.733
1036		GPU 64 COO min 3.980 max 4.110 mean 4.027	1110	Row-Gradient	H min 11.835 max 11.840 mean 11.838
1037		CSR min 4.320 max 4.490 mean 4.437	1111		GPU 64 COO min 2.240 max 2.390 mean 2.329
1038		H min 9.809 max 9.813 mean 9.811	1112		CSR min 4.570 max* 4.850 mean 4.807
1039	c-47.mat		1113		H min 10.742 max 10.752 mean 10.747
1040	Regular		1114	Column-Gradient	
1041		GPU 64 COO min 3.980 max* 4.080 mean 4.026	1115		GPU 64 COO min 2.010 max 2.050 mean 2.034
1042		CSR min 4.760 max 4.850 mean 4.812	1116		CSR min 4.570 max 4.760 mean 4.701
1043		H min 8.364 max 8.364 mean 8.364	1117		H min 11.872 max*11.881 mean 11.878
1044	Row-Premute		1118	Row-Column-Permute	
1045		GPU 64 COO min 3.880 max 4.010 mean 3.942	1119		GPU 64 COO min 2.000 max 2.040 mean 2.023
1046		CSR min 4.040 max 4.900 mean 4.807	1120		CSR min 0.770 max 4.780 mean 4.594
1047		H min 10.059 max 10.063 mean 10.061	1121		H min 11.835 max 11.840 mean 11.838
1048	Row-Gradient		1122		
1049		GPU 64 COO min 3.900 max 4.050 mean 3.976	1123	gen4.mat	
1050		CSR min 4.380 max 4.740 mean 4.630	1124	Regular	
1051		H min 10.201 max 10.228 mean 10.214	1125		GPU 64 COO min 4.880 max 4.980 mean 4.900
1052	Column-Gradient		1126		CSR min 10.020 max*11.300 mean 10.716
1053		GPU 64 COO min 3.860 max 3.990 mean 3.936	1127		H min 9.234 max 9.234 mean 9.234
1054		CSR min 4.350 max 4.610 mean 4.525	1128	Row-Premute	
1055		H min 11.204 max*11.241 mean 11.222	1129		GPU 64 COO min 4.860 max 4.930 mean 4.890
1056	Row-Column-Permute		1130		CSR min 0.330 max 11.200 mean 10.038
1057		GPU 64 COO min 3.890 max 4.020 mean 3.953	1131		H min 10.249 max 10.254 mean 10.252
1058		CSR min 4.490 max* 4.920 mean 4.840	1132	Row-Gradient	
1059		H min 10.058 max 10.063 mean 10.061	1133		GPU 64 COO min 4.860 max* 4.990 mean 4.908
1060	case9.mat		1134		CSR min 9.160 max 11.240 mean 10.435
1061	Regular		1135		H min 9.939 max 9.961 mean 9.947
1062		GPU 64 COO min 0.000 max 0.000 mean 0.000	1136	Column-Gradient	
1063		CSR min 0.000 max 0.000 mean 0.000	1137		GPU 64 COO min 4.780 max 4.880 mean 4.816

1138			CSR min 7.770 max 10.570 mean 9.407	1212	Row-Premute		
1139		H	min 10.851 max*10.876 mean 10.864	1213		GPU 64 COO min 4.420 max 4.520 mean 4.445	
1140	Row-Column-Permute			1214		CSR min 10.520 max 10.880 mean 10.696	
1141		GPU 64 COO min 4.850 max 4.950 mean 4.886		1215		H min 10.960 max*10.968 mean 10.963	
1142		CSR min 10.220 max 11.280 mean 10.748		1216	Row-Gradient		
1143		H min 10.250 max 10.255 mean 10.252		1217		GPU 64 COO min 4.570 max 4.690 mean 4.605	
1144	lp_fit2d.mtx			1218		CSR min 4.550 max 13.350 mean 12.479	
1145	Regular			1219		H min 9.508 max 9.527 mean 9.520	
1146		GPU 64 COO min 4.360 max* 4.640 mean 4.515		1220	Column-Gradient		
1147		CSR min 10.080 max 10.900 mean 10.491		1221		GPU 64 COO min 4.430 max 4.530 mean 4.461	
1148		H min 11.109 max 11.109 mean 11.109		1222		CSR min 10.250 max 10.940 mean 10.603	
1149	Row-Premute			1223		H min 10.934 max 10.945 mean 10.939	
1150		GPU 64 COO min 4.170 max 4.630 mean 4.476		1224	Row-Column-Permute		
1151		CSR min 0.910 max 10.910 mean 10.257		1225		GPU 64 COO min 4.420 max 4.520 mean 4.450	
1152		H min 11.098 max 11.104 mean 11.101		1226		CSR min 7.380 max 10.900 mean 10.598	
1153	Row-Gradient			1227		H min 10.959 max 10.967 mean 10.963	
1154		GPU 64 COO min 4.370 max 4.630 mean 4.529		1228	mult_dcop_01.mtx		
1155		CSR min 10.030 max 10.970 mean 10.624		1229	Regular		
1156		H min 11.109 max 11.109 mean 11.109		1230		GPU 64 COO min 3.420 max 3.630 mean 3.555	
1157	Column-Gradient			1231		CSR min 3.650 max 4.090 mean 3.814	
1158		GPU 64 COO min 4.250 max 4.640 mean 4.499		1232		H min 9.689 max 9.689 mean 9.689	
1159		CSR min 8.510 max*11.010 mean 10.505		1233	Row-Premute		
1160		H min 11.328 max*11.333 mean 11.331		1234		GPU 64 COO min 3.450 max 3.580 mean 3.521	
1161	Row-Column-Permute			1235		CSR min 3.610 max 4.150 mean 3.785	
1162		GPU 64 COO min 4.350 max 4.640 mean 4.511		1236		H min 10.738 max 10.742 mean 10.740	
1163		CSR min 10.040 max 10.790 mean 10.468		1237	Row-Gradient		
1164		H min 11.097 max 11.106 mean 11.101		1238		GPU 64 COO min 3.510 max* 3.660 mean 3.579	
1165	lp_osa_07.mtx			1239		CSR min 3.650 max 4.160 mean 3.806	
1166	Regular			1240		H min 10.576 max 10.585 mean 10.580	
1167		GPU 64 COO min 0.460 max* 3.640 mean 3.456		1241	Column-Gradient		
1168		CSR min 5.570 max* 8.530 mean 8.106		1242		GPU 64 COO min 3.460 max 3.650 mean 3.584	
1169		H min 8.412 max 8.412 mean 8.412		1243		CSR min 3.660 max* 4.240 mean 3.799	
1170	Row-Premute			1244		H min 10.826 max*10.842 mean 10.836	
1171		GPU 64 COO min 3.140 max 3.450 mean 3.367		1245	Row-Column-Permute		
1172		CSR min 7.600 max 8.070 mean 7.853		1246		GPU 64 COO min 3.470 max 3.580 mean 3.532	
1173		H min 9.255 max 9.258 mean 9.256		1247		CSR min 3.600 max 3.980 mean 3.743	
1174	Row-Gradient			1248		H min 10.738 max 10.742 mean 10.740	
1175		GPU 64 COO min 3.190 max 3.610 mean 3.509		1249	mult_dcop_02.mtx		
1176		CSR min 0.000 max 8.260 mean 7.597		1250	Regular		
1177		H min 8.583 max 8.678 mean 8.670		1251		GPU 64 COO min 3.390 max 3.660 mean 3.585	
1178	Column-Gradient			1252		CSR min 0.960 max 4.330 mean 4.162	
1179		GPU 64 COO min 3.330 max 3.500 mean 3.416		1253		H min 9.689 max 9.689 mean 9.689	
1180		CSR min 6.730 max 7.540 mean 7.199		1254	Row-Premute		
1181		H min 9.542 max* 9.604 mean 9.581		1255		GPU 64 COO min 3.310 max 3.600 mean 3.488	
1182	Row-Column-Permute			1256		CSR min 0.620 max 4.290 mean 4.132	
1183		GPU 64 COO min 3.290 max 3.430 mean 3.365		1257		H min 10.738 max 10.743 mean 10.740	
1184		CSR min 7.390 max 8.060 mean 7.832		1258	Row-Gradient		
1185		H min 9.255 max 9.258 mean 9.256		1259		GPU 64 COO min 3.310 max* 3.670 mean 3.593	
1186	Maragal_6.mtx			1260		CSR min 4.130 max* 4.430 mean 4.331	
1187	Regular			1261		H min 10.576 max 10.584 mean 10.580	
1188		GPU 64 COO min 4.160 max 4.310 mean 4.217		1262	Column-Gradient		
1189		CSR min 4.940 max 4.960 mean 4.956		1263		GPU 64 COO min 0.550 max 3.660 mean 3.486	
1190		H min 9.930 max 9.930 mean 9.930		1264		CSR min 3.890 max 4.410 mean 4.275	
1191	Row-Premute			1265		H min 10.831 max*10.843 mean 10.836	
1192		GPU 64 COO min 4.220 max 4.240 mean 4.225		1266	Row-Column-Permute		
1193		CSR min 4.750 max*13.040 mean 5.133		1267		GPU 64 COO min 3.470 max 3.590 mean 3.542	
1194		H min 10.776 max 10.778 mean 10.777		1268		CSR min 4.190 max 4.290 mean 4.242	
1195	Row-Gradient			1269		H min 10.738 max 10.742 mean 10.740	
1196		GPU 64 COO min 4.180 max* 4.450 mean 4.245		1270	mult_dcop_03.mtx		
1197		CSR min 4.880 max 4.940 mean 4.915		1271	Regular		
1198		H min 11.259 max*11.302 mean 11.281		1272		GPU 64 COO min 3.360 max* 3.660 mean 3.550	
1199	Column-Gradient			1273		CSR min 3.650 max 4.090 mean 3.813	
1200		GPU 64 COO min 4.200 max 4.250 mean 4.236		1274		H min 9.689 max 9.689 mean 9.689	
1201		CSR min 4.800 max 4.890 mean 4.859		1275	Row-Premute		
1202		H min 12.022 max 12.073 mean 12.051		1276		GPU 64 COO min 3.450 max 3.580 mean 3.521	
1203	Row-Column-Permute			1277		CSR min 3.610 max 4.160 mean 3.784	
1204		GPU 64 COO min 4.210 max 4.230 mean 4.222		1278		H min 10.738 max 10.743 mean 10.740	
1205		CSR min 4.860 max 4.890 mean 4.887		1279	Row-Gradient		
1206		H min 10.776 max 10.778 mean 10.778		1280		GPU 64 COO min 3.470 max 3.660 mean 3.572	
1207	mhd4800a.mtx			1281		CSR min 3.640 max 4.190 mean 3.809	
1208	Regular			1282		H min 10.572 max 10.584 mean 10.580	
1209		GPU 64 COO min 4.570 max* 4.710 mean 4.608		1283	Column-Gradient		
1210		CSR min 12.690 max*13.940 mean 13.369		1284		GPU 64 COO min 3.430 max 3.650 mean 3.562	
1211		H min 7.132 max 7.132 mean 7.132		1285		CSR min 3.670 max* 4.290 mean 3.793	

1286		H	min 10.828 max*10.840 mean 10.834	1360		GPU 64 COO min 4.540 max 4.940 mean 4.874
1287	Row-Column-Permute			1361		CSR min 6.280 max 6.520 mean 6.403
1288		GPU 64 COO min 3.370 max 3.610 mean 3.502		1362	H	min 10.042 max 10.047 mean 10.044
1289		CSR min 3.610 max 3.970 mean 3.744		1363	Row-Gradient	
1290		H	min 10.738 max 10.741 mean 10.740	1364		GPU 64 COO min 4.830 max 4.930 mean 4.875
1291	OPF_3754.mtx			1365		CSR min 5.790 max* 6.560 mean 6.289
1292	Regular			1366	H	min 9.675 max 9.706 mean 9.692
1293		GPU 64 COO min 4.700 max* 4.930 mean 4.842		1367	Column-Gradient	
1294		CSR min 6.230 max* 6.600 mean 6.411		1368		GPU 64 COO min 4.790 max* 4.960 mean 4.880
1295		H	min 8.393 max 8.393 mean 8.393	1369		CSR min 5.760 max 6.450 mean 6.204
1296	Row-Premute			1370	H	min 10.601 max*10.661 mean 10.626
1297		GPU 64 COO min 4.620 max 4.890 mean 4.787		1371	Row-Column-Permute	
1298		CSR min 5.780 max 6.310 mean 6.192		1372		GPU 64 COO min 4.330 max 4.950 mean 4.845
1299		H	min 11.265 max 11.272 mean 11.269	1373		CSR min 5.740 max 6.500 mean 6.375
1300	Row-Gradient			1374	H	min 10.041 max 10.046 mean 10.044
1301		GPU 64 COO min 4.570 max 4.870 mean 4.776		1375	TSOPF_RS_b39_c7.mtx	
1302		CSR min 5.770 max 6.510 mean 6.302		1376	Regular	
1303		H	min 10.464 max 10.473 mean 10.468	1377		GPU 64 COO min 4.300 max* 4.430 mean 4.364
1304	Column-Gradient			1378		CSR min 4.480 max 4.750 mean 4.716
1305		GPU 64 COO min 4.580 max 4.870 mean 4.756		1379	H	min 7.304 max 7.304 mean 7.304
1306		CSR min 5.630 max 6.180 mean 6.055		1380	Row-Premute	
1307		H	min 11.394 max*11.401 mean 11.397	1381		GPU 64 COO min 4.260 max 4.400 mean 4.353
1308	Row-Column-Permute			1382		CSR min 4.490 max 4.770 mean 4.734
1309		GPU 64 COO min 4.610 max 4.900 mean 4.780		1383	H	min 10.536 max 10.541 mean 10.539
1310		CSR min 5.010 max 6.300 mean 6.113		1384	Row-Gradient	
1311		H	min 11.268 max 11.272 mean 11.270	1385		GPU 64 COO min 3.970 max 4.420 mean 4.338
1312	OPF_6000.mtx			1386		CSR min 4.620 max* 4.820 mean 4.789
1313	Regular			1387	H	min 9.638 max 9.644 mean 9.641
1314		GPU 64 COO min 3.780 max* 3.920 mean 3.864		1388	Column-Gradient	
1315		CSR min 4.270 max 4.360 mean 4.332		1389		GPU 64 COO min 4.240 max 4.430 mean 4.368
1316		H	min 8.799 max 8.799 mean 8.799	1390		CSR min 4.710 max 4.770 mean 4.736
1317	Row-Premute			1391	H	min 11.129 max*11.222 mean 11.205
1318		GPU 64 COO min 3.770 max 3.870 mean 3.821		1392	Row-Column-Permute	
1319		CSR min 3.970 max*11.050 mean 4.439		1393		GPU 64 COO min 4.260 max 4.410 mean 4.359
1320		H	min 11.872 max 11.877 mean 11.875	1394		CSR min 4.660 max 4.760 mean 4.738
1321	Row-Gradient			1395	H	min 10.537 max 10.541 mean 10.539
1322		GPU 64 COO min 3.700 max 3.870 mean 3.795				
1323		CSR min 4.330 max 4.440 mean 4.403				
1324		H	min 11.109 max 11.116 mean 11.113			
1325	Column-Gradient			1396	11 FIJI	
1326		GPU 64 COO min 3.690 max 3.870 mean 3.804		1397	mult_dcop_03.mtx	
1327		CSR min 4.260 max 4.340 mean 4.308		1398	Regular	
1328		H	min 12.041 max*12.045 mean 12.043	1399		GPU 64 COO min 5.140 max* 5.140 mean 5.140
1329	Row-Column-Permute			1400		CSR min 10.340 max*10.390 mean 10.365
1330		GPU 64 COO min 3.780 max 3.860 mean 3.819		1401	H	min 9.689 max 9.689 mean 9.689
1331		CSR min 4.090 max 4.290 mean 4.259		1402	Row-Premute	
1332		H	min 11.873 max 11.877 mean 11.876	1403		GPU 64 COO min 4.970 max 4.990 mean 4.980
1333	shermanACb.mtx			1404		CSR min 9.420 max 9.430 mean 9.425
1334	Regular			1405	H	min 10.739 max 10.739 mean 10.739
1335		GPU 64 COO min 2.920 max* 3.140 mean 3.048		1406	Row-Gradient	
1336		CSR min 5.550 max 5.980 mean 5.803		1407		GPU 64 COO min 5.080 max 5.090 mean 5.085
1337		H	min 8.600 max 8.600 mean 8.600	1408		CSR min 9.720 max 10.300 mean 10.010
1338	Row-Premute			1409	H	min 10.579 max 10.582 mean 10.580
1339		GPU 64 COO min 2.760 max 3.020 mean 2.898		1410	Column-Gradient	
1340		CSR min 2.660 max 5.830 mean 5.632		1411		GPU 64 COO min 5.030 max 5.120 mean 5.075
1341		H	min 10.377 max 10.381 mean 10.379	1412		CSR min 9.330 max 9.770 mean 9.550
1342	Row-Gradient			1413	H	min 10.835 max*10.838 mean 10.836
1343		GPU 64 COO min 2.800 max 3.040 mean 2.944		1414	Row-Column-Permute	
1344		CSR min 5.330 max* 6.020 mean 5.742		1415		GPU 64 COO min 5.000 max 5.010 mean 5.005
1345		H	min 9.919 max 9.925 mean 9.922	1416		CSR min 7.580 max 9.460 mean 8.520
1346	Column-Gradient			1417	H	min 10.739 max 10.741 mean 10.740
1347		GPU 64 COO min 2.720 max 3.010 mean 2.926		1418	mult_dcop_03.mtx	
1348		CSR min 0.000 max 5.840 mean 5.513		1419	Regular	
1349		H	min 10.587 max*10.596 mean 10.591	1420		GPU 64 COO min 5.140 max* 5.140 mean 5.140
1350	Row-Column-Permute			1421		CSR min 10.340 max*10.390 mean 10.365
1351		GPU 64 COO min 2.780 max 3.030 mean 2.939		1422	H	min 9.689 max 9.689 mean 9.689
1352		CSR min 4.860 max 5.810 mean 5.667		1423	Row-Premute	
1353		H	min 10.376 max 10.382 mean 10.379	1424		GPU 64 COO min 4.970 max 4.990 mean 4.980
1354	TSOPF_FS_b9_c6.mtx			1425		CSR min 9.420 max 9.430 mean 9.425
1355	Regular			1426	H	min 10.739 max 10.739 mean 10.739
1356		GPU 64 COO min 0.000 max 0.000 mean 0.000		1427	Row-Gradient	
1357		CSR min 0.000 max 0.000 mean 0.000		1428		GPU 64 COO min 5.080 max 5.090 mean 5.085
1358		H	min 7.380 max 7.380 mean 7.380	1429		CSR min 9.720 max 10.300 mean 10.010
1359	Row-Premute			1430	H	min 10.579 max 10.582 mean 10.580

1431	Column-Gradient				1505		CSR min	6.360	max	7.450	mean	6.711
1432		GPU 64 COO min	5.030	max	5.120	mean	5.075					
1433		CSR min	9.330	max	9.770	mean	9.550					
1434		H	min	10.835	max*10.838	mean	10.836					
1435	Row-Column-Permute											
1436		GPU 64 COO min	5.000	max	5.010	mean	5.005					
1437		CSR min	7.580	max	9.460	mean	8.520					
1438		H	min	10.739	max	10.741	mean	10.740				
1439	mult_dcop_03.mtx											
1440	Regular											
1441		GPU 64 COO min	5.130	max*	5.220	mean	5.142					
1442		CSR min	7.250	max*	9.320	mean	7.722					
1443		H	min	9.689	max	9.689	mean	9.689				
1444	Row-Premute											
1445		GPU 64 COO min	4.980	max	5.030	mean	4.999					
1446		CSR min	6.460	max	8.470	mean	6.950					
1447		H	min	10.738	max	10.742	mean	10.740				
1448	Row-Gradient											
1449		GPU 64 COO min	5.070	max	5.140	mean	5.088					
1450		CSR min	6.780	max	8.700	mean	7.268					
1451		H	min	10.572	max	10.584	mean	10.580				
1452	Column-Gradient											
1453		GPU 64 COO min	4.980	max	5.030	mean	5.010					
1454		CSR min	6.390	max	7.640	mean	6.982					
1455		H	min	10.825	max*10.845	mean	10.836					
1456	Row-Column-Permute											
1457		GPU 64 COO min	4.990	max	5.010	mean	4.997					
1458		CSR min	6.300	max	7.160	mean	6.636					
1459		H	min	10.738	max	10.743	mean	10.740				
1460	mult_dcop_01.mtx											
1461	Regular											
1462		GPU 64 COO min	5.120	max*	5.140	mean	5.134					
1463		CSR min	6.990	max*	9.230	mean	7.546					
1464		H	min	9.689	max	9.689	mean	9.689				
1465	Row-Premute											
1466		GPU 64 COO min	4.990	max	5.020	mean	5.004					
1467		CSR min	6.370	max	7.220	mean	6.771					
1468		H	min	10.738	max	10.743	mean	10.740				
1469	Row-Gradient											
1470		GPU 64 COO min	5.060	max	5.100	mean	5.082					
1471		CSR min	6.730	max	7.720	mean	7.317					
1472		H	min	10.574	max	10.585	mean	10.580				
1473	Column-Gradient											
1474		GPU 64 COO min	4.980	max	5.100	mean	5.012					
1475		CSR min	6.580	max	7.510	mean	7.054					
1476		H	min	10.828	max*10.842	mean	10.835					
1477	Row-Column-Permute											
1478		GPU 64 COO min	4.970	max	5.000	mean	4.986					
1479		CSR min	6.390	max	7.050	mean	6.677					
1480		H	min	10.738	max	10.742	mean	10.740				
1481	mult_dcop_02.mtx											
1482	Regular											
1483		GPU 64 COO min	5.120	max	5.140	mean	5.133					
1484		CSR min	6.950	max	7.590	mean	7.336					
1485		H	min	9.689	max	9.689	mean	9.689				
1486	Row-Premute											
1487		GPU 64 COO min	4.970	max	4.990	mean	4.984					
1488		CSR min	6.440	max	7.110	mean	6.719					
1489		H	min	10.738	max	10.742	mean	10.740				
1490	Row-Gradient											
1491		GPU 64 COO min	5.070	max*	5.150	mean	5.086					
1492		CSR min	6.650	max*	7.930	mean	7.304					
1493		H	min	10.574	max	10.587	mean	10.580				
1494	Column-Gradient											
1495		GPU 64 COO min	4.980	max	5.040	mean	5.012					
1496		CSR min	6.520	max	7.650	mean	7.139					
1497		H	min	10.829	max*10.846	mean	10.836					
1498	Row-Column-Permute											
1499		GPU 64 COO min	4.970	max	5.050	mean	4.983					
1500		CSR min	6.440	max	7.380	mean	6.779					
1501		H	min	10.738	max	10.743	mean	10.740				
1502	lp_fit2d.mtx											
1503	Regular											
1504		GPU 64 COO min	3.960	max	3.960	mean	3.960					
1505												
1506												
1507	Row-Premute											
1508		GPU 64 COO min	3.950	max*	3.980	mean	3.953					
1509		CSR min	6.330	max	7.400	mean	6.661					
1510		H	min	11.098	max	11.104	mean	11.101				
1511	Row-Gradient											
1512		GPU 64 COO min	3.960	max	3.980	mean	3.961					
1513		CSR min	6.270	max*10.770	mean	7.017						
1514		H	min	11.109	max	11.109	mean	11.109				
1515	Column-Gradient											
1516		GPU 64 COO min	3.940	max	3.960	mean	3.950					
1517		CSR min	6.270	max	7.370	mean	6.696					
1518		H	min	11.329	max*11.334	mean	11.331					
1519	Row-Column-Permute											
1520		GPU 64 COO min	3.950	max	3.960	mean	3.952					
1521		CSR min	6.180	max	7.420	mean	6.641					
1522		H	min	11.098	max	11.105	mean	11.101				
1523	bloweya.mtx											
1524	Regular											
1525		GPU 64 COO min	0.000	max	0.000	mean	0.000					
1526		CSR min	0.000	max	0.000	mean	0.000					
1527		H	min	7.205	max	7.205	mean	7.205				
1528	Row-Premute											
1529		GPU 64 COO min	4.020	max	4.030	mean	4.023					
1530		CSR min	6.070	max	6.750	mean	6.340					
1531		H	min	11.025	max	11.031	mean	11.028				
1532	Row-Gradient											
1533		GPU 64 COO min	4.090	max*	4.160	mean	4.111					
1534		CSR min	5.980	max*	7.370	mean	6.678					
1535		H	min	10.295	max	10.304	mean	10.300				
1536	Column-Gradient											
1537		GPU 64 COO min	3.980	max	4.010	mean	3.995					
1538		CSR min	5.880	max	6.780	mean	6.295					
1539		H	min	10.881	max*10.887	mean	10.883					
1540	Row-Column-Permute											
1541		GPU 64 COO min	4.020	max	4.030	mean	4.023					
1542		CSR min	5.970	max	6.420	mean	6.183					
1543		H	min	11.025	max	11.033	mean	11.028				
1544	lp_osa_07.mtx											
1545	Regular											
1546		GPU 64 COO min	4.260	max*	4.270	mean	4.261					
1547		CSR min	6.440	max	7.640	mean	6.863					
1548		H	min	8.412	max	8.412	mean	8.412				
1549	Row-Premute											
1550		GPU 64 COO min	4.200	max	4.200	mean	4.200					
1551		CSR min	6.020	max	7.030	mean	6.418					
1552		H	min	9.255	max	9.257	mean	9.256				
1553	Row-Gradient											
1554		GPU 64 COO min	4.210	max	4.240	mean	4.226					
1555		CSR min	6.070	max*10.050	mean	6.498						
1556		H	min	8.607	max	8.678	mean	8.671				
1557	Column-Gradient											
1558		GPU 64 COO min	4.170	max	4.190	mean	4.180					
1559		CSR min	5.610	max	7.300	mean	5.988					
1560		H	min	9.534	max*	9.601	mean	9.585				
1561	Row-Column-Permute											
1562		GPU 64 COO min	4.190	max	4.190	mean	4.190					
1563		CSR min	6.070	max	7.000	mean	6.386					
1564		H	min	9.255	max	9.257	mean	9.256				
1565	ex19.mtx											
1566	Regular											
1567		GPU 64 COO min	6.140	max*	6.180	mean	6.159					
1568		CSR min	12.780	max*14.400	mean	13.328						
1569		H	min	8.228	max	8.228	mean	8.228				
1570	Row-Premute											
1571		GPU 64 COO min	5.820	max	5.850	mean	5.833					
1572		CSR min	9.870	max	11.070	mean	10.372					
1573		H	min	11.836	max	11.840	mean	11.838				
1574	Row-Gradient											
1575		GPU 64 COO min	6.070	max	6.120	mean	6.104					
1576		CSR min	11.290	max	12.760	mean	12.088					
1577		H	min	10.743	max							

1579		GPU 64 COO min 5.760 max 5.840 mean 5.813	1653		H min 7.380 max 7.380 mean 7.380
1580		CSR min 9.710 max 14.220 mean 10.376	1654	Row-Premute	
1581	H	min 11.873 max*11.882 mean 11.878	1655		GPU 64 COO min 4.130 max 4.170 mean 4.134
1582	Row-Column-Permute		1656		CSR min 6.180 max* 9.200 mean 6.796
1583		GPU 64 COO min 5.810 max 5.860 mean 5.838	1657	Row-Gradient	H min 10.041 max 10.046 mean 10.044
1584		CSR min 9.920 max 10.820 mean 10.240	1658		
1585	H	min 11.836 max 11.841 mean 11.838	1659		GPU 64 COO min 4.150 max* 4.220 mean 4.163
1586	brainpc2.mtx		1660		CSR min 6.410 max 7.500 mean 6.816
1587	Regular		1661		H min 9.682 max 9.706 mean 9.693
1588		GPU 64 COO min 0.000 max 0.000 mean 0.000	1662	Column-Gradient	
1589		CSR min 0.000 max 0.000 mean 0.000	1663		GPU 64 COO min 4.080 max 4.110 mean 4.096
1590	H	min 7.478 max 7.478 mean 7.478	1664		CSR min 6.020 max 7.220 mean 6.309
1591	Row-Premute		1665		H min 10.597 max*10.658 mean 10.631
1592		GPU 64 COO min 4.760 max 4.790 mean 4.773	1666	Row-Column-Permute	
1593		CSR min 6.930 max 7.780 mean 7.310	1667		GPU 64 COO min 4.120 max 4.140 mean 4.130
1594	H	min 9.810 max 9.813 mean 9.811	1668		CSR min 6.210 max 7.200 mean 6.609
1595	Row-Gradient		1669		H min 10.041 max 10.046 mean 10.044
1596		GPU 64 COO min 4.820 max* 4.840 mean 4.831	1670	TSOPF_FS_b9_c6.mtx	
1597		CSR min 7.220 max 8.290 mean 7.583	1671	Regular	
1598	H	min 9.721 max 9.725 mean 9.723	1672		GPU 64 COO min 0.000 max 0.000 mean 0.000
1599	Column-Gradient		1673		CSR min 0.000 max 0.000 mean 0.000
1600		GPU 64 COO min 4.760 max 4.820 mean 4.779	1674		H min 7.380 max 7.380 mean 7.380
1601		CSR min 6.870 max* 8.300 mean 7.393	1675	Row-Premute	
1602	H	min 10.368 max*10.373 mean 10.370	1676		GPU 64 COO min 4.120 max 4.140 mean 4.129
1603	Row-Column-Permute		1677		CSR min 6.170 max 7.160 mean 6.664
1604		GPU 64 COO min 4.750 max 4.780 mean 4.765	1678		H min 10.041 max 10.045 mean 10.043
1605		CSR min 6.940 max 7.580 mean 7.298	1679	Row-Gradient	
1606	H	min 9.809 max 9.814 mean 9.811	1680		GPU 64 COO min 4.150 max* 4.180 mean 4.162
1607	shermanAcB.mtx		1681		CSR min 6.420 max 7.360 mean 6.723
1608	Regular		1682		H min 9.682 max 9.706 mean 9.693
1609		GPU 64 COO min 4.090 max* 4.130 mean 4.112	1683	Column-Gradient	
1610		CSR min 6.320 max* 7.200 mean 6.779	1684		GPU 64 COO min 4.080 max 4.120 mean 4.096
1611	H	min 8.600 max 8.600 mean 8.600	1685		CSR min 5.880 max 7.090 mean 6.403
1612	Row-Premute		1686		H min 10.611 max*10.660 mean 10.637
1613		GPU 64 COO min 4.020 max 4.050 mean 4.036	1687	Row-Column-Permute	
1614		CSR min 5.670 max 6.460 mean 6.014	1688		GPU 64 COO min 4.130 max 4.140 mean 4.130
1615	H	min 10.376 max 10.382 mean 10.379	1689		CSR min 6.330 max* 7.390 mean 6.695
1616	Row-Gradient		1690		H min 10.042 max 10.047 mean 10.044
1617		GPU 64 COO min 4.050 max 4.100 mean 4.074	1691	OPF_6000.mtx	
1618		CSR min 5.580 max 6.420 mean 5.996	1692	Regular	
1619	H	min 9.918 max 9.924 mean 9.921	1693		GPU 64 COO min 7.270 max* 7.370 mean 7.293
1620	Column-Gradient		1694		CSR min 12.890 max*14.500 mean 13.566
1621		GPU 64 COO min 4.010 max 4.080 mean 4.033	1695		H min 8.799 max 8.799 mean 8.799
1622		CSR min 0.000 max 6.320 mean 5.527	1696	Row-Premute	
1623	H	min 10.543 max*10.595 mean 10.589	1697		GPU 64 COO min 6.640 max 6.720 mean 6.678
1624	Row-Column-Permute		1698		CSR min 9.680 max 11.600 mean 10.040
1625		GPU 64 COO min 4.020 max 4.050 mean 4.036	1699		H min 11.873 max 11.877 mean 11.875
1626		CSR min 5.670 max 6.510 mean 6.092	1700	Row-Gradient	
1627	H	min 10.377 max 10.381 mean 10.379	1701		GPU 64 COO min 7.090 max 7.140 mean 7.122
1628	cvxqp3.mtx		1702		CSR min 11.250 max 13.030 mean 12.142
1629	Regular		1703		H min 11.110 max 11.117 mean 11.114
1630		GPU 64 COO min 3.500 max* 3.540 mean 3.501	1704	Column-Gradient	
1631		CSR min 11.860 max*13.100 mean 12.694	1705		GPU 64 COO min 6.590 max 6.710 mean 6.644
1632	H	min 8.646 max 8.646 mean 8.646	1706		CSR min 9.400 max 13.140 mean 9.991
1633	Row-Premute		1707		H min 12.040 max*12.046 mean 12.043
1634		GPU 64 COO min 3.360 max 3.370 mean 3.365	1708	Row-Column-Permute	
1635		CSR min 6.210 max 7.610 mean 6.631	1709		GPU 64 COO min 6.640 max 6.710 mean 6.679
1636	H	min 11.027 max 11.032 mean 11.030	1710		CSR min 9.690 max 10.740 mean 10.050
1637	Row-Gradient		1711		H min 11.874 max 11.877 mean 11.875
1638		GPU 64 COO min 3.370 max 3.380 mean 3.376	1712	OPF_3754.mtx	
1639		CSR min 6.170 max 7.070 mean 6.499	1713	Regular	
1640	H	min 11.059 max 11.068 mean 11.064	1714		GPU 64 COO min 4.430 max* 4.450 mean 4.443
1641	Column-Gradient		1715		CSR min 9.710 max*13.000 mean 11.377
1642		GPU 64 COO min 3.350 max 3.390 mean 3.371	1716		H min 8.393 max 8.393 mean 8.393
1643		CSR min 6.150 max 7.180 mean 6.531	1717	Row-Premute	
1644	H	min 11.125 max*11.133 mean 11.130	1718		GPU 64 COO min 4.230 max 4.250 mean 4.240
1645	Row-Column-Permute		1719		CSR min 7.430 max 8.750 mean 7.986
1646		GPU 64 COO min 3.350 max 3.380 mean 3.364	1720		H min 11.266 max 11.272 mean 11.269
1647		CSR min 6.040 max 7.440 mean 6.603	1721	Row-Gradient	
1648	H	min 11.028 max 11.033 mean 11.030	1722		GPU 64 COO min 4.370 max 4.420 mean 4.382
1649	case9.mtx		1723		CSR min 8.160 max 9.470 mean 8.682
1650	Regular		1724		H min 10.462 max 10.473 mean 10.468
1651		GPU 64 COO min 0.000 max 0.000 mean 0.000	1725	Column-Gradient	
1652		CSR min 0.000 max 0.000 mean 0.000	1726		GPU 64 COO min 4.210 max 4.240 mean 4.227

1727		CSR min 7.160 max 8.080 mean 7.595	1801	Row-Premute	
1728		H min 11.394 max*11.401 mean 11.398	1802		GPU 64 COO min 10.340 max 10.430 mean 10.362
1729	Row-Column-Permute		1803		CSR min 12.880 max 13.340 mean 13.057
1730		GPU 64 COO min 4.230 max 4.250 mean 4.243	1804		H min 10.777 max 10.778 mean 10.777
1731		CSR min 7.230 max 8.940 mean 8.056	1805	Row-Gradient	
1732		H min 11.264 max 11.271 mean 11.269	1806		GPU 64 COO min 10.650 max*10.740 mean 10.688
1733	c-47.mtx		1807		CSR min 12.310 max 13.670 mean 12.562
1734	Regular		1808		H min 11.247 max 11.300 mean 11.281
1735		GPU 64 COO min 5.320 max* 5.340 mean 5.329	1809	Column-Gradient	
1736		CSR min 8.890 max* 9.590 mean 9.249	1810		GPU 64 COO min 10.340 max 10.440 mean 10.398
1737		H min 8.364 max 8.364 mean 8.364	1811		CSR min 9.480 max 10.110 mean 9.782
1738	Row-Premute		1812		H min 12.023 max*12.069 mean 12.047
1739		GPU 64 COO min 5.240 max 5.250 mean 5.241	1813	Row-Column-Permute	
1740		CSR min 7.790 max 8.890 mean 8.214	1814		GPU 64 COO min 10.330 max 10.380 mean 10.356
1741		H min 10.059 max 10.063 mean 10.061	1815		CSR min 12.840 max 13.530 mean 13.119
1742	Row-Gradient		1816		H min 10.776 max 10.778 mean 10.777
1743		GPU 64 COO min 5.230 max 5.260 mean 5.242	1817	aft01.mtx	
1744		CSR min 7.080 max 8.050 mean 7.673	1818	Regular	
1745		H min 10.206 max 10.226 mean 10.218	1819		GPU 64 COO min 3.680 max* 3.690 mean 3.688
1746	Column-Gradient		1820		CSR min 13.860 max*14.830 mean 14.560
1747		GPU 64 COO min 5.080 max 5.120 mean 5.105	1821		H min 7.811 max 7.811 mean 7.811
1748		CSR min 5.780 max 6.970 mean 6.359	1822	Row-Premute	
1749		H min 11.205 max*11.233 mean 11.222	1823		GPU 64 COO min 3.510 max 3.530 mean 3.513
1750	Row-Column-Permute		1824		CSR min 6.420 max 10.520 mean 7.265
1751		GPU 64 COO min 5.220 max 5.250 mean 5.227	1825		H min 11.161 max*11.170 mean 11.165
1752		CSR min 7.860 max 8.710 mean 8.247	1826	Row-Gradient	
1753		H min 10.059 max 10.064 mean 10.061	1827		GPU 64 COO min 3.630 max 3.670 mean 3.643
1754	mhd4800a.mtx		1828		CSR min 10.760 max 13.510 mean 12.199
1755	Regular		1829		H min 10.248 max 10.265 mean 10.258
1756		GPU 64 COO min 3.090 max* 3.100 mean 3.098	1830	Column-Gradient	
1757		CSR min 11.570 max*12.290 mean 12.092	1831		GPU 64 COO min 3.510 max 3.520 mean 3.519
1758		H min 7.132 max 7.132 mean 7.132	1832		CSR min 6.490 max 11.230 mean 7.645
1759	Row-Premute		1833		H min 11.112 max 11.121 mean 11.117
1760		GPU 64 COO min 3.020 max 3.020 mean 3.020	1834	Row-Column-Permute	
1761		CSR min 5.560 max 7.270 mean 6.007	1835		GPU 64 COO min 3.510 max 3.540 mean 3.515
1762		H min 10.959 max*10.968 mean 10.963	1836		CSR min 6.510 max 11.650 mean 7.311
1763	Row-Gradient		1837		H min 11.161 max 11.168 mean 11.165
1764		GPU 64 COO min 3.080 max 3.100 mean 3.088	1838	TSOPF_RS_b39_c7.mtx	
1765		CSR min 10.250 max 12.150 mean 11.340	1839	Regular	
1766		H min 9.509 max 9.528 mean 9.520	1840		GPU 64 COO min 5.970 max* 6.010 mean 5.988
1767	Column-Gradient		1841		CSR min 12.470 max*21.120 mean 13.816
1768		GPU 64 COO min 3.020 max 3.050 mean 3.026	1842		H min 7.304 max 7.304 mean 7.304
1769		CSR min 5.530 max 10.580 mean 6.432	1843	Row-Premute	
1770		H min 10.933 max 10.946 mean 10.939	1844		GPU 64 COO min 5.840 max 5.870 mean 5.856
1771	Row-Column-Permute		1845		CSR min 10.780 max 15.810 mean 11.425
1772		GPU 64 COO min 3.020 max 3.020 mean 3.020	1846		H min 10.537 max 10.540 mean 10.539
1773		CSR min 5.510 max 6.830 mean 6.136	1847	Row-Gradient	
1774		H min 10.959 max 10.967 mean 10.963	1848		GPU 64 COO min 5.950 max 6.000 mean 5.975
1775	gen4.mtx		1849		CSR min 11.520 max 17.250 mean 12.799
1776	Regular		1850		H min 9.638 max 9.646 mean 9.641
1777		GPU 64 COO min 3.300 max* 3.320 mean 3.308	1851	Column-Gradient	
1778		CSR min 5.250 max 6.340 mean 5.705	1852		GPU 64 COO min 5.790 max 5.860 mean 5.827
1779		H min 9.234 max 9.234 mean 9.234	1853		CSR min 10.500 max 14.080 mean 11.237
1780	Row-Premute		1854		H min 11.128 max*11.223 mean 11.209
1781		GPU 64 COO min 3.290 max 3.310 mean 3.299	1855	Row-Column-Permute	
1782		CSR min 5.190 max 7.420 mean 5.683	1856		GPU 64 COO min 5.850 max 5.870 mean 5.855
1783		H min 10.249 max 10.254 mean 10.252	1857		CSR min 10.790 max 15.250 mean 11.718
1784	Row-Gradient		1858		H min 10.537 max 10.541 mean 10.539
1785		GPU 64 COO min 3.300 max 3.310 mean 3.301	1859	mult_dcop_03.mtx	
1786		CSR min 5.370 max 6.310 mean 5.659	1860	Regular	
1787		H min 9.934 max 9.958 mean 9.948	1861		GPU 64 COO min 5.130 max* 5.220 mean 5.142
1788	Column-Gradient		1862		CSR min 7.250 max* 9.320 mean 7.722
1789		GPU 64 COO min 3.240 max 3.260 mean 3.249	1863		H min 9.689 max 9.689 mean 9.689
1790		CSR min 5.090 max* 8.660 mean 5.546	1864	Row-Premute	
1791		H min 10.853 max*10.873 mean 10.864	1865		GPU 64 COO min 4.980 max 5.030 mean 4.999
1792	Row-Column-Permute		1866		CSR min 6.460 max 8.470 mean 6.950
1793		GPU 64 COO min 3.290 max 3.320 mean 3.296	1867		H min 10.738 max 10.742 mean 10.740
1794		CSR min 5.190 max 7.550 mean 5.659	1868	Row-Gradient	
1795		H min 10.249 max 10.255 mean 10.252	1869		GPU 64 COO min 5.070 max 5.140 mean 5.088
1796	Maragal_6.mtx		1870		CSR min 6.780 max 8.700 mean 7.268
1797	Regular		1871		H min 10.572 max 10.584 mean 10.580
1798		GPU 64 COO min 10.580 max 10.620 mean 10.599	1872	Column-Gradient	
1799		CSR min 15.620 max*16.470 mean 15.832	1873		GPU 64 COO min 4.980 max 5.030 mean 5.010
1800		H min 9.930 max 9.930 mean 9.930	1874		CSR min 6.390 max 7.640 mean 6.982

```

1875           H           min 10.825 max*10.845 mean 10.836
1876 Row-Column-Permute
1877           GPU 64 COO min 4.990 max 5.010 mean 4.997
1878           CSR min 6.300 max 7.160 mean 6.636
1879           H           min 10.738 max 10.743 mean 10.740

```

REFERENCES

- 1881 [1] Hartwig Anzt, Terry Cojean, Chen Yen-Chen, Jack J. Dongarra, Goran
1882 Flegar, Pratik Nayak, Stanimire Tomov, Yuhsiang M. Tsai, and Weichung
1883 Wang. 2020. Load-balancing Sparse Matrix Vector Product Kernels
1884 on GPUs. *ACM Trans. Parallel Comput.* 7, 1 (2020), 2:1–2:26. <https://doi.org/10.1145/3380930>
- 1886 [2] Paolo D'Alberto, Chris Drome, and Ali Dasdan. 2012. Non-Parametric
1887 Methods Applied to the N-Sample Series Comparison. (2012).
1888 arXiv:stat.CO/1205.1880
- 1889 [3] Enver Kayaaslan, Cevdet Aykanat, and Bora Uçar. 2018. 1.5D Parallel
1890 Sparse Matrix-Vector Multiply. *SIAM J. Scientific Computing* 40, 1 (2018).
1891 <https://doi.org/10.1137/16M1105591>
- 1892 [4] Brian A. Page and Peter M. Kogge. 2018. Scalability of Hybrid Sparse
1893 Matrix Dense Vector (SpMV) Multiplication. In *2018 International Con-*
1894 *ference on High Performance Computing & Simulation, HPCS 2018, Or-*
1895 *leans, France, July 16-20, 2018*. IEEE, 406–414. <https://doi.org/10.1109/HPCS.2018.00072>